

Table of Contents

- 1 - Problem Statement
- 2 - Overview of Dataset
- 3 - Loading the Data
- 4 - Exploratory Data Analysis
 - 4.1 - Data Cleaning: Deduplication
- 5 - Text Preprocessing Using NLTK
- 6 - Train and Test Split of Data
- 7 - Logistic Regression Model
 - 7.1 - Function to find the optimal hyperparameter($C = 1/\lambda$) and error using K-fold cross-validation
 - 7.2 - Function to predict on Test Data
 - 7.3 - Function to check the behaviour of sparsity on weight vectors on increasing $C = 1/\lambda$ in L1 regularization
 - 7.4 - Multicollinearity Check (Perturbation Test)
 - 7.5 - Feature Importance
- 8 - Featurization Methods
 - 8.1 - Bag Of Words(unigram)
 - 8.2 - Bag Of Words(bigram)
 - 8.3 - TF-IDF(unigram)
 - 8.4 - TF-IDF(bigram)
 - 8.5 - Average Word2Vec
 - 8.6 - TF-IDF Weighted Word2Vec
- 9 - Conclusion

[1] Problem Statement :

- Time Based slicing to split Train Data(70%) and Test Data(30%).
- Applying Logistic Regression model to find the optimal hyperparameter(lambda) using 10 fold Cross Validation(both GridSearch and RandomSearch) in :
 - 1)Bag Of Words
 - 2)TF-IDF
 - 2)Average Word2Vec
 - 2)TF-IDF weighted Word2Vec
- Adding regularization(L1 and L2) to our logistic model.
- Checking the Accuracy and Sparsity of optimal weight vectors by increasing the value of lambda in L1 Regularization.
- Checking the Multicollinearity in weight vectors and finding the feature Importance.

[2] Overview of Dataset :

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>
(<https://www.kaggle.com/snap/amazon-fine-food-reviews>)

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454

Number of users: 256,059

Number of products: 74,258

Timespan: Oct 1999 - Oct 2012

Number of Attributes/Columns in data: 10

Attribute Information:

- 1.Id
- 2.ProductId - unique identifier for the product
- 3.UserId - unique identifier for the user
- 4.ProfileName
- 5.HelpfulnessNumerator - number of users who found the review helpful
- 6.HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
- 7.Score - rating between 1 and 5
- 8.Time - timestamp for the review
- 9.Summary - brief summary of the review
- 10.Text - text of the review

Objective: Given a review, determine whether the review is positive (Rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative?

[Ans] We could use the Score/Rating. A rating of 4 or 5 could be considered a positive review. A review of 1 or 2 could be considered negative. A review of 3 is neutral and ignored. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

[3] Loading the Data :

In order to load the data, we have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

In [1]:

```
#Importing the necessary Packages
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import time
import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from IPython.display import HTML
from collections import OrderedDict

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle
```

In [2]:

```
import pickle

#dumping an object to file object using dump method
def dumpfile(a,file_Name):
    fileObject = open(file_Name,"wb")
    pickle.dump(a,fileObject,protocol=2)
    fileObject.close()

#Loading an object from file object using load method
def loadfile(file_Name):
    fileObject = open(file_Name,"rb")
    b = pickle.load(fileObject)
    return b
```

In [3]:

```
%%HTML
<style type="text/css">
table.dataframe td, table.dataframe th {
    border: 2px  black solid !important;
}
</style>
```

In [4]:

```
# using the SQLite Table to read data.
con = sqlite3.connect('database.sqlite')

#filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 """, con)
```

In [7]:

```
# Give reviews with Score>3 a positive rating, and reviews with a score<3 a negative rating.
def partition(x):
    if x < 3:
        return 'negative'
    return 'positive'

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
```

In [8]:

```
print("Number of datapoints: ",filtered_data.shape[0])
print("Number of attributes/features: ",filtered_data.shape[1])
HTML(filtered_data.head().to_html(index=False))
```

Number of datapoints: 525814

Number of attributes/features: 10

Out[8]:

Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	Helpfulr
1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1
2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0
3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	1
4	B000UA0QIQ	A395BORC6FGVXV	Karl	3	3
5	B006K2ZZ7K	A1UQRSCLF8GW1T	Michael D. Bigham "M. Wassir"	0	0



[4] Exploratory Data Analysis :

[4.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

Deduplication 1:- As can be seen below the same user has multiple reviews of the with the same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that

ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8)

ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8)

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

In [9]:

```
display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", con)
```

```
HTML(display.head().to_html(index=False))
```

Out[9]:

Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator
78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2	2
138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2	2
138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2	2
73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2	2
155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2	2



In [10]:

```
#Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False, kind='quicksort', na_position='last')
```

In [11]:

```
#Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep=
'first', inplace=False)
final.shape
```

Out[11]:

(364173, 10)

Deduplication 2:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

In [12]:

```
display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)

HTML(display.head().to_html(index=False))
```

Out[12]:

Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator
64422	B000MIDROQ	A161DK06JJMCYF	J. E. Stephens "Jeanne"	3	1
44737	B001EQ55RW	A2V0I904FH7ABY	Ram	3	2

In [13]:

```
final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
print(final.shape)
```

(364171, 10)

Deduplication 3:- It was also seen that a same user has given different reviews for a same product at same time. I think it is normal for a user to give multiple reviews about a product, but that should be in different time. So, all those rows with same user giving multiple reviews for a same product at same time are considered as duplicate and hence dropped.

In [14]:

```
display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId= "A8891HVRDJAM6"
ORDER BY ProductID
""", con)

HTML(display.head().to_html(index=False))
```

Out[14]:

Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator
86221	B000084E6V	A8891HVRDJAM6	Marfaux "Marfaux"	33	33
86236	B000084E6V	A8891HVRDJAM6	Marfaux "Marfaux"	3	3

In [15]:

```
final=final.drop_duplicates(subset={"ProductId","UserId","ProfileName","Time"}, keep='first', inplace=False)
print(final.shape)
```

(363633, 10)

Deduplication 4:- It was also seen that in few rows with Ids from 150493 to 150529 contain reviews regarding books, not fine foods. So I think these should be also removed from the dataset. After looking at the productid column, it can be noticed that all the observations for fine foods start with B followed by numbers except for Ids from 150493 to 150529. I suppose the reviews for book 'Chicken soup for the soul' have gotten into the dataset mistakenly as they contain the words "chicken soup".

In [16]:

```
display = final[final.ProductId == "0006641040"]  
HTML(display.head().to_html(index=False))
```

Out[16]:

Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator
150524	0006641040	ACITT7DI6IDDL	shari zychinski	0	0
150506	0006641040	A2IW4PEEKO2R0U	Tracy	1	1
150507	0006641040	A1S4A3IQ2MU7V4	sally sue "sally sue"	1	1
150508	0006641040	AZGXZ2UUK6X	Catherine Hallberg " (Kate)"	1	1
150509	0006641040	A3CMRKGE0P909G	Teresa	3	4



In [17]:

```
final = final[final.ProductId != "0006641040"]
```

In [18]:

```
print("Percentage of data still remaining : ",(final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100)

#Before starting the next phase of preprocessing Lets see the number of entries left
print("Number of reviews left after Data Cleaning and Deduplication :")
print(final.shape)

#How many positive and negative reviews are present in our dataset?
final['Score'].value_counts()
```

```
Percentage of data still remaining : 69.14973735959865
Number of reviews left after Data Cleaning and Deduplication :
(363599, 10)
```

Out[18]:

```
positive    306566
negative     57033
Name: Score, dtype: int64
```

Observation:-

It is an imbalanced dataset as the number of positive reviews are way high in number than negative reviews.

[5] Text Preprocessing Using NLTK :

In the Preprocessing phase we do the following in the order below:-

1. Removal of HTML Tags
2. Removal of any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Removal of Stopwords
7. Finally Snowball Stemming the word

After which we collect the words used to describe positive and negative reviews

In [19]:

```
# find sentences containing HTML tags
import re
i=0;
for sent in final['Text'].values:
    if (len(re.findall('<.*?>', sent))):
        print(i)
        print(sent)
        break;
    i += 1;
```

10

I wanted a treat that was accepted and well liked for my rescue animals.
This is the only treat that is healthy and loved by all 4 legged beings in my home!
It does not contain sugar or grains or silly vegetables which virtually all treats contain. Dogs, cats and ferrets are carnivores they are not cattle to eat grain or rabbits to eat vegetables, and WHYYYY do companies add sugar,beet pulp or corn syrup to carnivore foods? It is dangerous and can cause the death of an animal with diabetes.
It is pretty easy to break into smaller pieces for cats and kittens with weak jaws and its wonderful to use as an aid to gain the trust of an abused dog as it will not cause stomach upset when given in common sense amounts.
I like that it goes a long way as it costs alot to heal and maintain and train abused and rescued dogs.
NO minus to this product other then the price,I can not afford to use it as much as I would like.

[5.1] Removal of html Tags :

In [20]:

```
#function to clean the word of any html-tags
def cleanhtml(sentence):
    cleanr = re.compile('<.*?>')
    cleantext = re.sub(cleanr, ' ', sentence)
    return cleantext
```

```
cleanhtml("<br />This is the only treat that is healthy and loved by all 4 legged beings in my home!<br />It does not contain sugar or grains")
```

Out[20]:

```
' This is the only treat that is healthy and loved by all 4 legged beings in my home! It does not contain sugar or grains'
```

[5.2] Removal of Punctuations and unnecessary characters :

In [21]:

```
#function to clean the word of any punctuation or special characters
def cleanpunc(sentence):
    cleaned = re.sub(r'[?|!|\'|"|#|@|~|%|*]',r'',sentence)
    cleaned = re.sub(r'[.,|)|(|\|/]',r' ',cleaned)
    return cleaned

cleanpunc("WHYYYY do companies add sugar,beet pulp or corn syrup to carnivore foods?")
```

Out[21]:

```
'WHYYYY do companies add sugar beet pulp or corn syrup to carnivore foods'
```

[5.3] StopWords :

In [22]:

```
import nltk
nltk.download('punkt')
nltk.download('stopwords')

stop = set(stopwords.words('english')) #set of stopwords
print(stop)
print("*****")
print("No. of stop words: ",len(stop))
```

```
[nltk_data] Downloading package punkt to /home/jovyan/nltk_data...
[nltk_data] Unzipping tokenizers/punkt.zip.
[nltk_data] Downloading package stopwords to /home/jovyan/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip.
{'they', 'having', 'by', 'other', 'you', 'from', 'that', 'which', 'he', 'n
eedn't', 'shan', 'mightn't', 'd', 'me', 'wouldn', 'has', 'yourselves', 'he
re', 'both', 'shan't', 'doesn't', 'can', 'is', 'ain', 'into', 'or', 'sam
e', 't', 'again', 'you're', 'are', 'mustn', 'under', 'further', 'hers', 'a
s', 'himself', 'hadn't', 'few', 'have', 'below', 'hasn', 'off', 'until',
'my', 'she', 's', 'should', 'once', 'its', 'nor', 'where', 'haven', 'mor
e', 'how', 'too', 'just', 'were', 'don't', 'of', 'hasn't', 'before', 'coul
dn't', 'm', 'shouldn't', 'not', 'hadn', 'her', 'ourselves', 'it's', 'our',
'but', 'after', 'did', 'haven't', 'herself', 'am', 'themselves', 'be',
'a', 'wasn't', 'ma', 'weren't', 'own', 'these', 'yours', 'wasn', 'itself',
'an', 'your', 'shouldn', 'weren', 'being', 'out', 'to', 'their', 'him', 'y
ourself', 'won't', 'you'd', 've', 'only', 'aren', 'during', 'couldn', 'tho
se', 'isn', 're', 'against', 'all', 'you've', 'at', 'some', 'aren't', 'an
d', 'above', 'the', 'about', 'wouldn't', 'over', 'because', 'was', 'does',
'will', 'what', 'no', 'who', 'most', 'very', 'up', 'isn't', 'with', 'now',
'should've', 'won', 'y', 'needn', 'you'll', 'his', 'why', 'theirs', 'who
m', 'don', 'it', 'doing', 'each', 'had', 'do', 'we', 'ours', 'while', 'i',
'myself', 'there', 'such', 'o', 'any', 'been', 'she's', 'when', 'between',
'll', 'doesn', 'didn't', 'mustn't', 'than', 'through', 'on', 'then', 'the
m', 'so', 'didn', 'that'll', 'down', 'this', 'for', 'mightn', 'if', 'in'}
*****
No. of stop words: 179
```

In [23]:

```
exceptions = ["aren't", "mightn", "wasn", "hadn", "don't", "against", "hadn't", "shan", "were  
n't", "didn", "don", "hasn't", \  
             "hasn", "shouldn", "didn't", "wouldn", "wasn't", "needn't", "shouldn't", "would  
n't", "aren", "isn't", "doesn't", \  
             "nor", "not", "needn", "couldn't", "mightn't", "mustn", "mustn't", "ain", "sha  
n't", "haven", "won't", "couldn", "isn", \  
             "weren", "haven't", "no", "haven't"]  
  
new_stop = []  
for i in stop:  
    if i not in exceptions:  
        new_stop.append(i)  
  
print("No. of stop words after removing exceptions: ", len(new_stop))
```

No. of stop words after removing exceptions: 140

[5.4] Stemming :

In [24]:

```
sno = nltk.stem.SnowballStemmer('english') #initialising the snowball stemmer  
  
print("Original word: beautiful" + "|" + "Stem word: " + sno.stem('beautiful'))  
print("Original word: beauty" + "|" + "Stem word: " + sno.stem('beauty'))  
print("Original word: loved" + "|" + "Stem word: " + sno.stem('loved'))  
print("Original word: loving" + "|" + "Stem word: " + sno.stem('loving'))
```

```
Original word: beautiful|Stem word: beauti  
Original word: beauty|Stem word: beauti  
Original word: loved|Stem word: love  
Original word: loving|Stem word: love
```

Observation:-

We can see words like "beautiful" and "beauty" have their stem as "beauti", "loved" and "loving" have their stem as "love".

Hence it helps in reducing the dimensions by taking the root stem of words.

[5.5] Implementing the preprocessing steps one by one on all the reviews of dataset :

In [22]:

```
i=0
str1=' '
final_string=[]
all_positive_words=[] # store words from +ve reviews here
all_negative_words=[] # store words from -ve reviews here.
s=''

for sent in final['Text'].values:
    filtered_sentence=[]
    #print(sent);
    sent=cleanhtml(sent) # remove HTML tags
    for w in sent.split():
        for cleaned_words in cleanpunc(w).split():
            if((cleaned_words.isalpha()) & (len(cleaned_words)>2)):
                if(cleaned_words.lower() not in new_stop):
                    s=(sno.stem(cleaned_words.lower())).encode('utf8')
                    filtered_sentence.append(s)
                    if (final['Score'].values)[i] == 'positive':
                        all_positive_words.append(s) #list of all words used to describ
e positive reviews
                    if(final['Score'].values)[i] == 'negative':
                        all_negative_words.append(s) #list of all words used to describ
e negative reviews reviews
                    else:
                        continue
                else:
                    continue
    str1 = b" ".join(filtered_sentence) #final string of cleaned words
    #print("*****")

    final_string.append(str1)
    i+=1
```

In [23]:

```
from nltk.probability import FreqDist
pdist = FreqDist(all_positive_words)
top_positive = pdist.most_common(20)
print("Top 20 Positive words ocuring frequently in reviews:")
top_positive
```

Top 20 Positive words ocuring frequently in reviews:

Out[23]:

```
[(b'not', 146568),
 (b'like', 139160),
 (b'tast', 128865),
 (b'good', 112601),
 (b'flavor', 109329),
 (b'love', 107172),
 (b'use', 103792),
 (b'great', 103670),
 (b'one', 96529),
 (b'product', 90912),
 (b'tri', 86683),
 (b'tea', 83699),
 (b'coffe', 78763),
 (b'make', 75004),
 (b'get', 71996),
 (b'food', 64539),
 (b'would', 55477),
 (b'time', 55184),
 (b'buy', 54137),
 (b'realli', 52657)]
```


In [24]:

```
ndist = FreqDist(all_negative_words)
top_negative = ndist.most_common(20)
print("Top 20 Negative words ocuring frequently in reviews:")
top_negative
```

Top 20 Negative words ocuring frequently in reviews:

Out[24]:

```
[(b'not', 54325),
 (b'tast', 34534),
 (b'like', 32271),
 (b'product', 28181),
 (b'one', 20544),
 (b'flavor', 19520),
 (b'would', 17947),
 (b'tri', 17718),
 (b'use', 15280),
 (b'good', 15024),
 (b'coffe', 14700),
 (b'get', 13775),
 (b'buy', 13742),
 (b'order', 12862),
 (b'food', 12720),
 (b'dont', 11865),
 (b'tea', 11646),
 (b'even', 11068),
 (b'box', 10833),
 (b'amazon', 10067)]
```

[5.6] Adding a new column of CleanedText which displays the data after pre-processing of the review :

In [27]:

```
final['CleanedText']=final_string
final['CleanedText']=final['CleanedText'].str.decode("utf-8")
```

In [26]:

```
final[['Text', 'CleanedText']].head(10)
```

Out[26]:

	Text	CleanedText
476617	This product by Arche...	product archer farm b...
22621	Our dogs just love th...	dog love saw pet stor...
22620	My dogs loves this ch...	dog love chicken prod...
284375	This book is easy to ...	book easi read ingred...
157850	I have been feeding m...	feed greyhound treat ...
157849	This is one product t...	one product welsh ter...
157833	This is the ONLY dog ...	dog treat lhasa apso ...
157832	These liver treas are...	liver trea phenomen r...
157837	This was the only tre...	treat dog like obedi ...
157831	No waste , even if sh...	wast even day goe hun...

[5.7] Using SQLite Table to load data after preprocessing of reviews :

In []:

```
# store final result into an SQLite table for future.
conn = sqlite3.connect('final.sqlite')
c=conn.cursor()
conn.text_factory = str
final.to_sql('Reviews', conn, schema=None, if_exists='replace', index=True, index_label=None, chunksize=None, dtype=None)
```

In [26]:

```
# using the SQLite Table to read data.
conn = sqlite3.connect('final.sqlite')

final = pd.read_sql_query(""" SELECT * FROM Reviews """,conn)
```

In [27]:

```
#Listing out the number of positive and negative reviews  
final = final.reset_index(drop=True)  
final['Score'].value_counts()
```

Out[27]:

```
positive    306566  
negative     57033  
Name: Score, dtype: int64
```

In [28]:

```
(final['Score'].value_counts()/len(final['Score']))*100
```

Out[28]:

```
positive     84.314313  
negative     15.685687  
Name: Score, dtype: float64
```

[6] Train and Test Split of Data :

Sorting the data by Time :

In [29]:

```
final=final.sort_values('Time', axis=0, ascending=True, inplace=False, kind='quicksort', na_position='last')
final.head()
```

Out[29]:

	index	Id	ProductId	UserId	ProfileName	HelpfulnessNun
387	417839	451856	B00004CXX9	AIUWLEQ1ADEG5	Elizabeth Medina	0
293	346055	374359	B00004CI84	A344SMIA5JECGM	Vincent P. Ross	1
386	417838	451855	B00004CXX9	AJH6LUC1UT1ON	The Phantom of the Opera	0
209	346116	374422	B00004CI84	A1048CYU0OV4O8	Judy L. Eans	2
271	346041	374343	B00004CI84	A1B2IZU1JLZA6	Wes	19

Considering negative reviews to be 0 and positive reviews to be 1 :

In [31]:

```
def reviews(x):
    if x == "positive":
        return 1
    else:
        return 0

final['Score'] = final['Score'].map(reviews)
final.head()
```

Out[31]:

	index	Id	ProductId	UserId	ProfileName	Helpfulness
387	417839	451856	B00004CXX9	AIUWLEQ1ADEG5	Elizabeth Medina	0
293	346055	374359	B00004CI84	A344SMIA5JECGM	Vincent P. Ross	1
386	417838	451855	B00004CXX9	AJH6LUC1UT1ON	The Phantom of the Opera	0
209	346116	374422	B00004CI84	A1048CYU0OV4O8	Judy L. Eans	2
271	346041	374343	B00004CI84	A1B2IZU1JLZA6	Wes	19

Time Based Slicing :

- Diving the data to Train set(first 70% ie older data) and Test Set(last 30% ie recent data)

In [32]:

```
from sklearn.model_selection import train_test_split

X = final["CleanedText"].values
y = final["Score"].values
X_train,X_test,y_train,y_test = train_test_split(X, y, test_size = 0.3,shuffle = False)
```

In [33]:

```
print("Shape of X_train: ",X_train.shape)
print("Shape of y_train: ",y_train.shape)
print("Shape of X_test: ",X_test.shape)
print("Shape of y_test: ",y_test.shape)
```

```
Shape of X_train: (254519,)
Shape of y_train: (254519,)
Shape of X_test: (109080,)
Shape of y_test: (109080,)
```

In [34]:

```
dumpfile(X,"X")
dumpfile(y,"y")
dumpfile(X_train,"X_train")
dumpfile(y_train,"y_train")
dumpfile(X_test,"X_test")
dumpfile(y_test,"y_test")
```

In [4]:

```
X = loadfile("X")
y = loadfile("y")
X_train = loadfile("X_train")
y_train = loadfile("y_train")
X_test = loadfile("X_test")
y_test = loadfile("y_test")
```

[7] Logistic Regression :

Grid Search Crossvalidation:

[7.1] Function to find the optimal hyperparameter($C = 1/\lambda$) and error using K-fold cross-validation :

- Taking C between range 0.0001 and 1000.
- Performing 10 fold cross validation on Train Data
- Finding the optimal C
- Plotting between CV error/CV Accuracy and $C(1/\lambda)$

In []:

```
from sklearn.model_selection import TimeSeriesSplit
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import cross_val_score as cv
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.grid_search import GridSearchCV
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score
from sklearn.metrics import roc_curve, auc
import warnings
warnings.filterwarnings('ignore')
```

```

def LR_gridTrain(X_train,y_train,penalty):

    C_values = [0.0001,0.001, 0.01, 0.1, 1, 5 , 10, 50, 100, 1000]
    param_grid = dict(C = C_values)

    model = LogisticRegression(penalty = penalty)
    grid = GridSearchCV(model, param_grid, cv=10, scoring='accuracy')
    grid.fit(X_train, y_train)

    grid_mean_scores = [i.mean_validation_score for i in grid.grid_scores_]

    #Misclassification error
    MSE = [1 - x for x in grid_mean_scores]

    #Finding the optimal K
    optimal_C = grid.best_params_

    best_accuracy = np.round(grid.best_score_ * 100,3)

    print("\n\033[1mOptimal C:\033[0m ", optimal_C)
    print("\n\033[1mCrossValidation Error:\033[0m {}".format(np.round(min(MSE),3)))
    print("\n\033[1mCrossValidation Accuracy:\033[0m {} %\n\n".format(best_accuracy))

    plt.figure(figsize=(20,6))
    plt.style.use('fivethirtyeight')
    plt.subplot(121)
    plt.plot(C_values,MSE, 'r-o')
    for xy in zip(C_values, np.round(MSE,3)):
        plt.annotate('(%s %s)' % xy, xy = xy, textcoords = 'data')
    plt.title("CV Error vs C = 1/λ ")
    plt.xlabel("C = 1/λ")
    plt.ylabel("CV Error")
    plt.grid(True)

    plt.subplot(122)
    plt.plot(C_values,grid_mean_scores, 'g-o')
    for xy in zip(C_values, np.round(grid_mean_scores,3)):
        plt.annotate('(%s %s)' % xy, xy = xy, textcoords = 'data')
    plt.title("CV Accuracy vs C = 1/λ")
    plt.xlabel("C = 1/λ")
    plt.ylabel("CV Accuracy")
    plt.grid(True)
    plt.show()

    print("\n\033[1mCV Error for each value of C:\033[0m ",np.round(MSE,3))
    print("\n\033[1mCV Accuracy for each value of C:\033[0m ",np.round(grid_mean_scores
,3))

```

RandomSearch Crossvalidation: To find a better hyperparameter value in fewer number of experiments. Here the hyperparameter ($C = 1/\lambda$) is defined as a distribution rather than list of values. Here, I use exponential distribution.

In [6]:

```
from sklearn.grid_search import RandomizedSearchCV
from scipy.stats import expon
import warnings
warnings.filterwarnings('ignore')

def LR_randomTrain(X_train,y_train,penalty):

    C_dist = expon(scale = 2)
    param_grid = dict(C = C_dist)

    model = LogisticRegression(penalty = penalty)
    random_grid = RandomizedSearchCV(model, param_grid, cv=10, scoring='accuracy')
    random_grid.fit(X_train, y_train)

    #Finding the optimal K
    optimal_C = random_grid.best_params_

    #Cv Accuracy
    best_accuracy = np.round(random_grid.best_score_ * 100,3)

    #Cv Error
    MSE = np.round((1 -random_grid.best_score_),3)

    print("\n\033[1mOptimal C:\033[0m ", optimal_C)
    print("\n\033[1mCrossValidation Error:\033[0m {}".format(MSE))
    print("\n\033[1mCrossValidation Accuracy:\033[0m {} %\n\n".format(best_accuracy))
```

[7.2] Function to predict on Test Data :

- Plotting the Confusion matrix
- Plotting the ROC/AUC Curve
- Finding Accuracy, Precision, Recall and F1 Score on Test Data

True Negative : Number of datapoints with class label "negative" correctly classified as "negative".

False Positive : Number of datapoints with class label "negative" misclassified as "positive".

False Negative : Number of datapoints with class label "positive" misclassified as "negative".

True Positive: Number of datapoints with class label "positive" correctly classified as "positive".

Precision : $\text{True Positive} / (\text{True Positive} + \text{False Positive})$

Recall : $\text{True Positive} / (\text{True Positive} + \text{False Negative})$

In [7]:

```
def LR_Test(X_train,X_test,y_train,y_test,penalty,optimal_C):

    optimal_model = LogisticRegression(C = optimal_C, penalty = penalty)
    optimal_model.fit(X_train, y_train)
    y_pred = optimal_model.predict(X_test)

    ##-----Confusion Matrix and Performance metrics-----##
    accuracy = accuracy_score(y_test,y_pred) * 100
    precision = precision_score(y_test,y_pred,average= 'macro')
    recall = recall_score(y_test,y_pred,average= 'macro')
    f1= f1_score(y_test,y_pred,average= 'macro')
    MSE = (1 - (accuracy/100))
    cm = confusion_matrix(y_test, y_pred)
    tn, fp, fn, tp = confusion_matrix(y_test, y_pred).ravel()
    cm_df = pd.DataFrame(cm,
                        index = ['negative','positive'],
                        columns = ['negative','positive'])
    sns.heatmap(cm_df, annot=True)
    plt.title('Confusion Matrix')
    plt.ylabel('Actual Label')
    plt.xlabel('Predicted Label')
    plt.show()

    print(cm)
    print("\n\033[1mTest Error :\033[0m {}".format(np.round(MSE,3)))
    print("\033[1mTest Accuracy :\033[0m {} %".format(np.round(accuracy,3)))
    print("\033[1mTrue Negative :\033[0m {}".format(tn))
    print("\033[1mFalse Positive :\033[0m {}".format(fp))
    print("\033[1mFalse Negative :\033[0m {}".format(fn))
    print("\033[1mTrue Positive :\033[0m {}".format(tp))
    print("\33[1mPrecision Score :\033[0m {}".format(np.round(precision,3)))
    print("\33[1mRecall Score :\033[0m {}".format(np.round(recall,3)))
    print("\33[1mF1 Score :\033[0m {}".format(np.round(f1,3)))

    print("\n\n")

    ##-----ROC Curve-----##
    fpr,tpr,thresholds = roc_curve(y_test,y_pred)
    roc_auc = auc(fpr,tpr)

    plt.figure(figsize=(8,6))
    plt.style.use('fivethirtyeight')
    plt.plot(fpr,tpr,'b',label="AUC = {}".format(roc_auc))
    plt.plot([0,1],[0,1],'r--')
    plt.ylabel('True Positive Rate')
    plt.xlabel('False Positive Rate')
    plt.title("Receiver Operating Characteristic")
    plt.legend()
    plt.grid(True)
    plt.show()
```

[7.3] Function to check the behaviour of sparsity on weight vectors on increasing $C = 1/\lambda$ in L1 regularization :

In [8]:

```
def sparsity_check(X_train,X_test,y_train,y_test):
    C_values = [10,1,0.1,0.01,0.001,0.0001]

    for i in C_values :
        start = time.clock()
        print("\n\033[1mSparsity and Accuarcy when C = {}\033[0m".format(i))
        model = LogisticRegression(penalty = 'l1',C = i)
        model.fit(X_train,y_train)
        y_pred = model.predict(X_test)
        accuracy = accuracy_score(y_test,y_pred) * 100
        train_error = 1 - model.score(X_train,y_train)
        test_error = 1 - model.score(X_test,y_test)

        print("Number of non-zero weighhts: ",np.count_nonzero(model.coef_))
        print("Train Error: ",np.round(train_error,3))
        print("Test Error: ",np.round(test_error,3))
        print("Test Accuracy :",np.round(accuracy_score(y_test,y_pred),5))
        print("Run Time :{}\033[0m sec".format(time.clock() - start))
```

[7.4] Multicollinearity Check (Pertubation Test) :

- Finding the weight vector with X_train
- Adding a small random noise with X_train and finding the new weight vecor by fitting to logistic model
- Setting a threshold of 30%,if the difference between orginal weight vector and new weight vector is more than 30%,then said to be multicollinear

In [9]:

```
from scipy.sparse import find

def multicollinear_check(X_train,X_test,y_train,y_test,optimal_C,threshold = 30):

    print("\033[1m-----BEFORE PERTUBATION TEST-----\033[0m\n")

    model = LogisticRegression(C = optimal_C, penalty = 'l2')
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    accuracy = accuracy_score(y_test,y_pred) * 100
    weights = model.coef_
    print("\033[1mSample Weights:\033[0m ",pd.DataFrame(weights))
    print("\033[1mSize of weight vector:\033[0m ",weights.size)
    print("\033[1mNon zero weights:\033[0m ",np.count_nonzero(model.coef_))
    print("\033[1mTest Accuracy :\033[0m {} %".format(np.round(accuracy,3)))

    #-----Adding a small random nosie(epsilon) to weight vector-----
    -----#
    X_trainnew = X_train
    #contains the row indices, column indices and values of the nonzero matrixes
    row,column,values = find(X_trainnew)

    np.random.seed(0)
    noise = np.random.normal(loc = 0, scale = 0.01, size = (find(X_trainnew)[0].size,))
    .astype(np.float16)
    X_trainnew[row,column] = noise + X_trainnew[row,column]

    print("\n\033[1m-----AFTER PERTUBATION TEST-----\033[0m\n")

    model_new = LogisticRegression(C = optimal_C, penalty = 'l2')
    model_new.fit(X_trainnew, y_train)
    y_prednew = model_new.predict(X_test)
    accuracy_new = accuracy_score(y_test,y_prednew) * 100
    weights_new = model_new.coef_
    print("\033[1mSample Weights:\033[0m ",pd.DataFrame(weights_new))
    print("\033[1mSize of weight vector:\033[0m ",weights_new.size)
    print("\033[1mNon zero weights:\033[0m ",np.count_nonzero(model_new.coef_))
    print("\033[1mTest Accuracy :\033[0m {} %".format(np.round(accuracy_new,3)))

    weight_difference = abs((weights - weights_new)/weights) * 100
    total = (weight_difference > threshold).sum()
    print("\n\n\033[1mNumber of features with weights changing greater than 30% :\033[0m",total)
    print("\n\n\033[1mFollowing are the {} features that are multicollinear\033[0m".format(total))
    size = np.shape(weight_difference)[1]
    for i in range(size):
        if weight_difference[0][i] > threshold:
            print(i, end = ' ')
```

[7.4] Feature Importance :

- Only finding the important features for each class in BOW and TFIDF

In [10]:

```
def feature_imortance(X_train,X_test,y_train,y_test,optimal_C,vectorizer, n = 25):

    model = LogisticRegression(C = optimal_C, penalty = 'l2')
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

    #-----Feature Importance-----#
    class_labels = model.classes_
    feature_names = vectorizer.get_feature_names()
    top_negative = sorted(zip(model.coef_[0], feature_names))[:n]
    top_positive= sorted(zip(model.coef_[0], feature_names))[-n:]

    print("\n\n\033[1m-----Top {} Negative Words with high Importance-----
-\033[0m".format(n))
    neg_featureimp_df = pd.DataFrame(top_negative, columns=['Coefficient Factor','Features'])
    print(neg_featureimp_df.to_string(index=False))

    print("\n\n\033[1m-----Top {} Positive Words with high Importance-----
\033[0m".format(n))
    pos_featureimp_df = pd.DataFrame(top_positive, columns=['Coefficient Factor','Features'],)
    print(pos_featureimp_df.to_string(index=False))
```

[8] Featurization Methods :

[8.1] Bag Of Words(unigram) :

In [46]:

```
%%time
bow_unigram = CountVectorizer(min_df = 0.0005)
X_train_bowuni = bow_unigram.fit_transform(X_train)
print("Type of Count Vectorizer: ",type(X_train_bowuni))
print("The shape of text BOW vectorizer: ", X_train_bowuni.get_shape())
print("Number of unique word: ", X_train_bowuni.get_shape()[1])
```

```
Type of Count Vectorizer: <class 'scipy.sparse.csr.csr_matrix'>
The shape of text BOW vectorizer: (254519, 3878)
Number of unique word: 3878
CPU times: user 11.6 s, sys: 32 ms, total: 11.6 s
Wall time: 11.6 s
```

In [47]:

```
%%time
X_test_bowuni = bow_unigram.transform(X_test)
print("The shape of text BOW vectorizer: ", X_test_bowuni.get_shape())
print("Number of unique word: ", X_test_bowuni.get_shape()[1])
```

The shape of text BOW vectorizer: (109080, 3878)
Number of unique word: 3878
CPU times: user 5.53 s, sys: 4 ms, total: 5.54 s
Wall time: 5.54 s

In [48]:

```
dumpfile(X_train_bowuni,"X_train_bowuni")
dumpfile(X_test_bowuni,"X_test_bowuni")
```

In [11]:

```
X_train_bowuni = loadfile("X_train_bowuni")
X_test_bowuni = loadfile("X_test_bowuni")
```

In [12]:

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler(with_mean=False)
X_train_bowuni_std = sc.fit_transform(X_train_bowuni)
```

In [13]:

```
X_test_bowuni_std = sc.transform(X_test_bowuni)
```

In [14]:

```
print("Shape of Training Data: ",X_train_bowuni_std.shape)
print("Shape of Test Data: ",X_test_bowuni_std.shape)
```

Shape of Training Data: (254519, 3878)
Shape of Test Data: (109080, 3878)

[8.1.1] Using GridSearch CV :

L2 Regularization :


In [15]:

```
%%time
if __name__ == "__main__":
    LR_gridTrain(X_train_bowuni_std, y_train, penalty = 'l2')
```

Optimal C: {'C': 0.001}

CrossValidation Error: 0.076

CrossValidation Accuracy: 92.384 %



CV Error for each value of C: [0.079 0.076 0.077 0.077 0.077 0.077 0.077 0.077 0.077 0.077]


CV Accuracy for each value of C: [0.921 0.924 0.923 0.923 0.923 0.923 0.923 0.923 0.923 0.923]

CPU times: user 9min 25s, sys: 2.56 s, total: 9min 27s

Wall time: 9min 27s


In [78]:

```
%%time
if __name__ == "__main__":
    LR_Test(X_train_bowuni_std, X_test_bowuni_std, y_train, y_test, 'l2', 0.001)
```

```
[[12473  6574]
 [ 2623 87410]]
```

Test Error : 0.084
 Test Accuracy : 91.569 %
 True Negative : 12473
 False Positive : 6574
 False Negative : 2623
 True Positive : 87410
 Precision Score : 0.878
 Recall Score : 0.813
 F1 Score : 0.84



CPU times: user 3.68 s, sys: 8 ms, total: 3.69 s
Wall time: 3.33 s

L1 Regularization :


In [75]:

```
%%time
if __name__ == "__main__":
    LR_gridTrain(X_train_bowuni_std, y_train, penalty = 'l1')
```

Optimal C: {'C': 0.1}

CrossValidation Error: 0.077

CrossValidation Accuracy: 92.325 %




CV Error for each value of C: [0.149 0.099 0.077 0.077 0.077 0.077 0.077
0.077 0.077 0.077]

CV Accuracy for each value of C: [0.851 0.901 0.923 0.923 0.923 0.923 0.923
0.923 0.923 0.923]

CPU times: user 3min 52s, sys: 2.01 s, total: 3min 54s
Wall time: 3min 54s


In [79]:

```
%%time
if __name__ == "__main__":
    LR_Test(X_train_bowuni_std, X_test_bowuni_std, y_train, y_test, 'l1', 0.1)
```



```
[[12611  6436]
 [ 2783 87250]]
```

Test Error : 0.085
Test Accuracy : 91.548 %
True Negative : 12611
False Positive : 6436
False Negative : 2783
True Positive : 87250
Precision Score : 0.875
Recall Score : 0.816
F1 Score : 0.841



CPU times: user 3.9 s, sys: 28 ms, total: 3.92 s
Wall time: 3.61 s

[8.1.2] Sparsity And Accuracy on Weight Vectors(L1 Regularization) :

In [17]:

```
if __name__ == "__main__":  
    sparsity_check(X_train_bowuni_std, X_test_bowuni_std ,y_train ,y_test)
```

Sparsity and Accuarcy when C = 10
Number of non-zero weighhts: 3876
Train Error: 0.069
Test Error: 0.084
Test Accuracy : 0.91569
Run Time :2.9262080000000004 sec

Sparsity and Accuarcy when C = 1
Number of non-zero weighhts: 3850
Train Error: 0.069
Test Error: 0.084
Test Accuracy : 0.91566
Run Time :2.668768 sec

Sparsity and Accuarcy when C = 0.1
Number of non-zero weighhts: 3714
Train Error: 0.069
Test Error: 0.085
Test Accuracy : 0.91547
Run Time :2.7338950000000004 sec

Sparsity and Accuarcy when C = 0.01
Number of non-zero weighhts: 2539
Train Error: 0.072
Test Error: 0.085
Test Accuracy : 0.91451
Run Time :2.7871330000000007 sec

Sparsity and Accuarcy when C = 0.001
Number of non-zero weighhts: 349
Train Error: 0.097
Test Error: 0.109
Test Accuracy : 0.891
Run Time :1.5401930000000004 sec

Sparsity and Accuarcy when C = 0.0001
Number of non-zero weighhts: 11
Train Error: 0.149
Test Error: 0.174
Test Accuracy : 0.82554
Run Time :1.144057 sec

Observation : Here $C = 1/\lambda$, we can see as C decreases(λ increases)

- Sparsity Increases(Number of non zero elements decreases)
- Error increases and Performance accuarcy drops(model starts underfitting)
- Run Time is also fast as sparsity increases

[8.1.3] Using RandomSearch CV :

L2 Regularization :

In [96]:

```
%%time  
if __name__ == "__main__":  
    LR_randomTrain(X_train_bowuni_std, y_train, penalty = 'l2')
```

Optimal C: {'C': 0.07806145232697875}

CrossValidation Error: 0.077


CrossValidation Accuracy: 92.299 %

CPU times: user 10min 47s, sys: 624 ms, total: 10min 48s

Wall time: 10min 48s


In [97]:

```
%%time
if __name__ == "__main__":
    LR_Test(X_train_bowuni_std, X_test_bowuni_std, y_train, y_test, 'l2', 0.078)
```

```
[[12692  6355]  
 [ 2845 87188]]
```

Test Error : 0.084
Test Accuracy : 91.566 %
True Negative : 12692
False Positive : 6355
False Negative : 2845
True Positive : 87188
Precision Score : 0.874
Recall Score : 0.817
F1 Score : 0.842



CPU times: user 5.82 s, sys: 12 ms, total: 5.83 s
Wall time: 5.46 s

L1 Regularization :

In [94]:

```
%%time  
if __name__ == "__main__":  
    LR_randomTrain(X_train_bowuni_std, y_train, penalty = 'l1')
```

Optimal C: {'C': 0.3521393406501542}


CrossValidation Error: 0.077

CrossValidation Accuracy: 92.3 %

CPU times: user 4min 31s, sys: 1.7 s, total: 4min 33s
Wall time: 4min 33s


In [95]:

```
%%time
if __name__ == "__main__":
    LR_Test(X_train_bowuni_std, X_test_bowuni_std, y_train, y_test, 'l1', 0.352)
```



```
[[12670  6377]
 [ 2826 87207]]
```

Test Error : 0.084
Test Accuracy : 91.563 %
True Negative : 12670
False Positive : 6377
False Negative : 2826
True Positive : 87207
Precision Score : 0.875
Recall Score : 0.817
F1 Score : 0.842



CPU times: user 3.96 s, sys: 40 ms, total: 4 s
Wall time: 3.67 s

[8.1.4] MultiCollinearity :

In [49]:

```
if __name__ == "__main__":  
    multicollinear_check(X_train_bowuni_std,X_test_bowuni_std,y_train,y_test,0.001)
```

-----BEFORE PERTUBATION TEST-----

```
Sample Weights:      0      1      2      3      4  
5      6      \  
0 0.000651 0.054345 0.037677 0.000138 0.004119 -0.008378 0.007636  
  
      7      8      9      ...      3868      3869      3870 \  
0 -0.012763 0.01727 0.022878 ...      0.012175 0.011759 0.004692  
  
      3871      3872      3873      3874      3875      3876      3877  
0 -0.017501 0.013488 0.025814 0.025312 0.010537 0.016239 -0.002453
```

[1 rows x 3878 columns]
Size of weight vector: 3878
Non zero weights: 3878
Test Accuracy : 91.574 %

-----AFTER PERTUBATION TEST-----

```
Sample Weights:      0      1      2      3      4      5  
6      \  
0 0.000654 0.054352 0.0377 0.000136 0.004148 -0.008354 0.007643  
  
      7      8      9      ...      3868      3869      3870 \  
0 -0.012724 0.017277 0.022884 ...      0.012187 0.011759 0.004678  
  
      3871      3872      3873      3874      3875      3876      3877  
0 -0.017495 0.013392 0.025791 0.025295 0.010553 0.016242 -0.00246
```

[1 rows x 3878 columns]
Size of weight vector: 3878
Non zero weights: 3878
Test Accuracy : 91.575 %

Number of features with weights changing greater than 30% : 11

Following are the 11 features that are multicollinear
11 479 1054 1626 1743 1797 1834 3007 3271 3495 3508

[8.1.5] FeatureImportance :

In [53]:

```
if __name__ == "__main__":  
    feature_importance(X_train_bowuni_std,X_test_bowuni_std,y_train,y_test,0.001,bow_uni  
gram)
```

-----Top 25 Negative Words with high Importance-----

Coefficient	Factor	Features
-0.426936		not
-0.281176		disappoint
-0.209263		worst
-0.171752		aw
-0.171402		terribl
-0.161840		horribl
-0.157774		tast
-0.157565		return
-0.147626		money
-0.144013		unfortun
-0.136553		thought
-0.132869		stale
-0.129352		would
-0.128290		bland
-0.125914		threw
-0.121448		didnt
-0.121260		wast
-0.118510		even
-0.112495		weak
-0.112317		bad
-0.110978		mayb
-0.110584		noth
-0.110162		hope
-0.108740		product
-0.106904		yuck

-----Top 25 Positive Words with high Importance-----

Coefficient	Factor	Features
0.148925		keep
0.153824		satisfi
0.154850		beat
0.157583		glad
0.161687		addict
0.162531		enjoy
0.163918		yummi
0.167071		awesom
0.170322		smooth
0.176150		find
0.183192		thank
0.193714		wonder
0.201873		happi
0.202285		easi
0.210035		tasti
0.217571		amaz
0.230436		favorit
0.248848		nice
0.297098		excel
0.336582		perfect
0.363671		good
0.392887		delici
0.439635		best
0.443822		love
0.604014		great

[8.2] Bag Of Words(bigram) :

In [13]:

```
%%time
bow_bigram = CountVectorizer(dtype='float', ngram_range=(1, 2), min_df = 0.0005)
X_train_bowbi = bow_bigram.fit_transform(X_train)
print("Type of Count Vectorizer: ", type(X_train_bowbi))
print("The shape of text BOW vectorizer: ", X_train_bowbi.get_shape())
print("Number of unique word: ", X_train_bowbi.get_shape()[1])
```

```
Type of Count Vectorizer: <class 'scipy.sparse.csr.csr_matrix'>
The shape of text BOW vectorizer: (254519, 10711)
Number of unique word: 10711
CPU times: user 38.4 s, sys: 444 ms, total: 38.8 s
Wall time: 38.8 s
```

In [14]:

```
%%time
X_test_bowbi = bow_bigram.transform(X_test)
print("The shape of text BOW vectorizer: ", X_test_bowbi.get_shape())
print("Number of unique word: ", X_test_bowbi.get_shape()[1])
```

```
The shape of text BOW vectorizer: (109080, 10711)
Number of unique word: 10711
CPU times: user 11.3 s, sys: 36 ms, total: 11.4 s
Wall time: 11.4 s
```

In [15]:

```
dumpfile(X_train_bowbi, "X_train_bowbi")
dumpfile(X_test_bowbi, "X_test_bowbi")
```

In [137]:

```
X_train_bowbi = loadfile("X_train_bowbi")
X_test_bowbi = loadfile("X_test_bowbi")
```

In [138]:

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler(with_mean=False)
X_train_bowbi_std = sc.fit_transform(X_train_bowbi)
```

In [139]:

```
X_test_bowbi_std = sc.transform(X_test_bowbi)
```

In [140]:

```
print("Shape of Training Data: ", X_train_bowbi_std.shape)
print("Shape of Test Data: ", X_test_bowbi_std.shape)
```

```
Shape of Training Data: (254519, 10711)
Shape of Test Data: (109080, 10711)
```

[8.2.1] Using GridSearch CV :

L2 Regularization :


In [124]:

```
%%time
if __name__ == "__main__":
    LR_gridTrain(X_train_bowbi_std, y_train, penalty = 'l2')
```

Optimal C: {'C': 0.0001}

CrossValidation Error: 0.064

CrossValidation Accuracy: 93.631 %



CV Error for each value of C: [0.064 0.065 0.068 0.069 0.069 0.069 0.069
0.069 0.069 0.069]


CV Accuracy for each value of C: [0.936 0.935 0.932 0.931 0.931 0.931 0.931
0.931 0.931 0.931]

CPU times: user 1h 8min 23s, sys: 35.5 s, total: 1h 8min 59s

Wall time: 31min


In [126]:

```
%%time
if __name__ == "__main__":
    LR_Test(X_train_bowbi_std, X_test_bowbi_std, y_train, y_test, 'l2', 0.0001)
```



```
[[13262  5785]
 [ 1846 88187]]
```

Test Error : 0.07
Test Accuracy : 93.004 %
True Negative : 13262
False Positive : 5785
False Negative : 1846
True Positive : 88187
Precision Score : 0.908
Recall Score : 0.838
F1 Score : 0.868



CPU times: user 7.15 s, sys: 48 ms, total: 7.2 s
Wall time: 3.9 s

L1 Regularization :


In [128]:

```
%%time
if __name__ == "__main__":
    LR_gridTrain(X_train_bowbi_std, y_train, penalty = 'l1')
```

Optimal C: {'C': 0.01}

CrossValidation Error: 0.065

CrossValidation Accuracy: 93.542 %



CV Error for each value of C: [0.149 0.09 0.065 0.067 0.069 0.069 0.069 0.069 0.069 0.069]


CV Accuracy for each value of C: [0.851 0.91 0.935 0.933 0.931 0.931 0.931 0.931 0.931 0.931]

CPU times: user 9min 59s, sys: 3.12 s, total: 10min 2s

Wall time: 10min 2s


In [130]:

```
%%time
if __name__ == "__main__":
    LR_Test(X_train_bowbi_std, X_test_bowbi_std, y_train, y_test, 'l1', 0.01)
```



```
[[13501  5546]
 [ 2047 87986]]
```

Test Error : 0.07
Test Accuracy : 93.039 %
True Negative : 13501
False Positive : 5546
False Negative : 2047
True Positive : 87986
Precision Score : 0.905
Recall Score : 0.843
F1 Score : 0.87



CPU times: user 4.68 s, sys: 48 ms, total: 4.73 s
Wall time: 4.39 s

[8.2.2] Sparsity And Accuracy on Weight Vectors(L1 Regularization) :

In [23]:

```
if __name__ == "__main__":  
    sparsity_check(X_train_bowbi_std, X_test_bowbi_std ,y_train ,y_test)
```

Sparsity and Accuarcy when C = 10
Number of non-zero weighhts: 10700
Train Error: 0.046
Test Error: 0.072
Test Accuracy : 0.92789
Run Time :8.503957 sec

Sparsity and Accuarcy when C = 1
Number of non-zero weighhts: 10635
Train Error: 0.046
Test Error: 0.072
Test Accuracy : 0.92802
Run Time :8.91921 sec

Sparsity and Accuarcy when C = 0.1
Number of non-zero weighhts: 9917
Train Error: 0.047
Test Error: 0.07
Test Accuracy : 0.92952
Run Time :6.138832000000001 sec

Sparsity and Accuarcy when C = 0.01
Number of non-zero weighhts: 5251
Train Error: 0.055
Test Error: 0.07
Test Accuracy : 0.93036
Run Time :3.1172979999999998 sec

Sparsity and Accuarcy when C = 0.001
Number of non-zero weighhts: 473
Train Error: 0.087
Test Error: 0.097
Test Accuracy : 0.90337
Run Time :1.7636930000000035 sec

Sparsity and Accuarcy when C = 0.0001
Number of non-zero weighhts: 12
Train Error: 0.149
Test Error: 0.174
Test Accuracy : 0.82583
Run Time :1.483823000000001 sec

Observation : Here $C = 1/\lambda$, we can see as C decreases(λ increases)

- Sparsity Increases(Number of non zero elements decreases)
- Error increases and Performance accuarcy drops(model starts underfitting)
- Run Time is also fast as sparsity increases

[8.2.3] Using RandomSearch CV :

L2 Regularization :

In [129]:

```
%%time  
if __name__ == "__main__":  
    LR_randomTrain(X_train_bowbi_std, y_train, penalty = 'l2')
```

Optimal C: {'C': 0.11828601922296544}

CrossValidation Error: 0.069


CrossValidation Accuracy: 93.079 %

CPU times: user 1h 28min 14s, sys: 39 s, total: 1h 28min 53s

Wall time: 39min 33s


In [131]:

```
%%time
if __name__ == "__main__":
    LR_Test(X_train_bowbi_std, X_test_bowbi_std, y_train, y_test, 'l2', 0.118)
```



```
[[14218  4829]
 [ 3031 87002]]
```

Test Error : 0.072
Test Accuracy : 92.794 %
True Negative : 14218
False Positive : 4829
False Negative : 3031
True Positive : 87002
Precision Score : 0.886
Recall Score : 0.856
F1 Score : 0.87



CPU times: user 35.4 s, sys: 228 ms, total: 35.6 s
Wall time: 17.4 s

L1 Regularization :

In [132]:

```
%%time  
if __name__ == "__main__":  
    LR_randomTrain(X_train_bowbi_std, y_train, penalty = 'l1')
```

Optimal C: {'C': 0.14447500879175584}


CrossValidation Error: 0.067

CrossValidation Accuracy: 93.257 %

CPU times: user 12min 6s, sys: 3.4 s, total: 12min 9s
Wall time: 12min 9s


In [133]:

```
%%time
if __name__ == "__main__":
    LR_Test(X_train_bowbi_std, X_test_bowbi_std, y_train, y_test, 'l1', 0.144)
```



```
[[14179  4868]
 [ 2855 87178]]
```

Test Error : 0.071
Test Accuracy : 92.92 %
True Negative : 14179
False Positive : 4868
False Negative : 2855
True Positive : 87178
Precision Score : 0.89
Recall Score : 0.856
F1 Score : 0.872



CPU times: user 7.41 s, sys: 48 ms, total: 7.46 s
Wall time: 7.09 s

[8.2.4] MultiCollinearity :

In [141]:

```
if __name__ == "__main__":  
    multicollinear_check(X_train_bowbi_std, X_test_bowbi_std, y_train, y_test,0.0001)
```

-----BEFORE PERTUBATION TEST-----

Sample Weights:	0	1	2	3	4	5
6 \						
0	0.001976	0.02532	0.017236	-0.007174	0.000338	0.021763
7		8	9	...	10701	10702
0	-0.003079	0.020889	0.013498	...	-0.009104	0.006508
	10704	10705	10706	10707	10708	10709
0	0.015023	0.00658	0.00078	0.006959	0.006199	0.004551
						10710
0						-0.001347

[1 rows x 10711 columns]

Size of weight vector: 10711

Non zero weights: 10711

Test Accuracy : 93.004 %

-----AFTER PERTUBATION TEST-----

Sample Weights:	0	1	2	3	4	5
6 \						
0	0.001975	0.025319	0.017242	-0.007172	0.000351	0.021761
7		8	9	...	10701	10702
0	-0.003079	0.020887	0.013507	...	-0.00911	0.006512
	10704	10705	10706	10707	10708	10709
0	0.01502	0.006573	0.00078	0.00695	0.006194	0.004554
						10710
0						-0.001337

[1 rows x 10711 columns]

Size of weight vector: 10711

Non zero weights: 10711

Test Accuracy : 93.005 %

Number of features with weights changing greater than 30% : 15

Following are the 15 features that are multicollinear

1158 1393 2595 3472 3842 3947 3980 4384 4702 6678 6916 7027 8211 8426 1023
3

[8.2.5] FeatureImportance :

In [143]:

```
if __name__ == "__main__":  
    feature_importance(X_train_bowbi_std,X_test_bowbi_std,y_train,y_test,0.0001,bow_bigram)
```

-----Top 25 Negative Words with high Importance-----

Coefficient	Factor	Features
-0.187487		disappoint
-0.112822		worst
-0.106119		not worth
-0.102862		not buy
-0.097046		not good
-0.094704	not	recommend
-0.094104		aw
-0.092579		terribl
-0.092061		not
-0.087579		horribl
-0.084748		unfortun
-0.081491		return
-0.078822		two star
-0.077773		stale
-0.076762		wont buy
-0.073707		threw
-0.073405		bland
-0.072333		weak
-0.072200	wast	money
-0.071753		bad
-0.068186		disgust
-0.067072		mayb
-0.065819		sorri
-0.062725	never	buy
-0.060210		tasteless

-----Top 25 Positive Words with high Importance-----

Coefficient	Factor	Features
0.070663		keep
0.071755	great	product
0.073626		satisfi
0.077443	tast	great
0.080625		addict
0.081148		enjoy
0.084625		yummi
0.084781		awesom
0.087047		easi
0.092139	not	disappoint
0.093123		amaz
0.094706		wonder
0.096731		thank
0.103235	high	recommend
0.103558		nice
0.104545		happi
0.107154		favorit
0.116067		tasti
0.126167		excel
0.132739		perfect
0.166082		good
0.173295		delici
0.174904		best
0.209069		love
0.233951		great

[8.3] TF-IDF(unigram) :

In [38]:

```
%%time
tfidf_unigram = TfidfVectorizer(dtype='float',min_df = 0.0005)
X_train_tfidfuni = tfidf_unigram.fit_transform(X_train)
print("Type of Count Vectorizer: ",type(X_train_tfidfuni))
print("The shape of text TFIDF vectorizer: ", X_train_tfidfuni.get_shape())
print("Number of unique word: ", X_train_tfidfuni.get_shape()[1])
```

```
Type of Count Vectorizer: <class 'scipy.sparse.csr.csr_matrix'>
The shape of text TFIDF vectorizer: (254519, 3878)
Number of unique word: 3878
CPU times: user 12 s, sys: 88 ms, total: 12.1 s
Wall time: 12.1 s
```

In [39]:

```
%%time
X_test_tfidfuni = tfidf_unigram.transform(X_test)
print("The shape of text TFIDF vectorizer: ", X_test_tfidfuni.get_shape())
print("Number of unique word: ", X_test_tfidfuni.get_shape()[1])
```

```
The shape of text TFIDF vectorizer: (109080, 3878)
Number of unique word: 3878
CPU times: user 5.6 s, sys: 16 ms, total: 5.62 s
Wall time: 5.62 s
```

In [40]:

```
dumpfile(X_train_tfidfuni,"X_train_tfidfuni")
dumpfile(X_test_tfidfuni,"X_test_tfidfuni")
```

In [144]:

```
X_train_tfidfuni = loadfile("X_train_tfidfuni")
X_test_tfidfuni = loadfile("X_test_tfidfuni")
```

In [145]:

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler(with_mean=False)
X_train_tfidfuni_std = sc.fit_transform(X_train_tfidfuni)
```

In [146]:

```
X_test_tfidfuni_std = sc.transform(X_test_tfidfuni)
```

In [147]:

```
print("Shape of Training Data: ",X_train_tfidfuni_std.shape)
print("Shape of Test Data: ",X_test_tfidfuni_std.shape)
```

```
Shape of Training Data: (254519, 3878)
Shape of Test Data: (109080, 3878)
```

[8.3.1] Using GridSearch CV :

L2 Regularization :


In [141]:

```
%%time
if __name__ == "__main__":
    LR_gridTrain(X_train_tfidfuni_std, y_train, penalty = 'l2')
```

Optimal C: {'C': 0.001}

CrossValidation Error: 0.074

CrossValidation Accuracy: 92.617 %



CV Error for each value of C: [0.079 0.074 0.074 0.074 0.074 0.074 0.074 0.074 0.074 0.074]


CV Accuracy for each value of C: [0.921 0.926 0.926 0.926 0.926 0.926 0.926 0.926 0.926 0.926]

CPU times: user 6min 43s, sys: 436 ms, total: 6min 44s

Wall time: 6min 43s


In [145]:

```
%%time
if __name__ == "__main__":
    LR_Test(X_train_tfidfuni_std, X_test_tfidfuni_std, y_train, y_test, 'l2', 0.001)
```



```
[[12913  6134]
 [ 2785 87248]]
```

Test Error : 0.082
Test Accuracy : 91.823 %
True Negative : 12913
False Positive : 6134
False Negative : 2785
True Positive : 87248
Precision Score : 0.878
Recall Score : 0.824
F1 Score : 0.847



CPU times: user 3.66 s, sys: 4 ms, total: 3.66 s
Wall time: 3.3 s

L1 Regularization :


In [142]:

```
%%time
if __name__ == "__main__":
    LR_gridTrain(X_train_tfidfuni_std, y_train, penalty = 'l1')
```

Optimal C: {'C': 0.1}

CrossValidation Error: 0.074

CrossValidation Accuracy: 92.619 %




CV Error for each value of C: [0.149 0.093 0.074 0.074 0.074 0.074 0.074 0.074 0.074 0.074]

CV Accuracy for each value of C: [0.851 0.907 0.926 0.926 0.926 0.926 0.926 0.926 0.926 0.926]

CPU times: user 7min 48s, sys: 1.98 s, total: 7min 50s
Wall time: 7min 50s

In [144]:

```
%%time
if __name__ == "__main__":
    LR_Test(X_train_tfidfuni_std, X_test_tfidfuni_std, y_train, y_test, 'l1', 0.1)
```



```
[[13120  5927]
 [ 2917 87116]]
```

Test Error : 0.081

Test Accuracy : 91.892 %

True Negative : 13120

False Positive : 5927


False Negative : 2917

True Positive : 87116

Precision Score : 0.877

Recall Score : 0.828

F1 Score : 0.85



CPU times: user 7.09 s, sys: 20 ms, total: 7.11 s
Wall time: 6.72 s

[8.3.2] Sparsity And Accuracy on Weight Vectors(L1 Regularization) :

In [28]:

```
if __name__ == "__main__":  
    sparsity_check(X_train_tfidfuni_std, X_test_tfidfuni_std ,y_train ,y_test)
```

Sparsity and Accuarcy when C = 10
Number of non-zero weighhts: 3875
Train Error: 0.066
Test Error: 0.081
Test Accuracy : 0.91856
Run Time :6.766615999999999 sec

Sparsity and Accuarcy when C = 1
Number of non-zero weighhts: 3853
Train Error: 0.066
Test Error: 0.081
Test Accuracy : 0.91868
Run Time :6.6327430000000005 sec

Sparsity and Accuarcy when C = 0.1
Number of non-zero weighhts: 3747
Train Error: 0.066
Test Error: 0.081
Test Accuracy : 0.91891
Run Time :6.4040860000000007 sec

Sparsity and Accuarcy when C = 0.01
Number of non-zero weighhts: 2682
Train Error: 0.068
Test Error: 0.081
Test Accuracy : 0.91913
Run Time :4.1066780000000002 sec

Sparsity and Accuarcy when C = 0.001
Number of non-zero weighhts: 409
Train Error: 0.09
Test Error: 0.102
Test Accuracy : 0.89846
Run Time :2.20008 sec

Sparsity and Accuarcy when C = 0.0001
Number of non-zero weighhts: 12
Train Error: 0.149
Test Error: 0.174
Test Accuracy : 0.8256
Run Time :1.1756900000000003 sec

Observation : Here $C = 1/\lambda$, we can see as C decreases(λ increases)

- Sparsity Increases(Number of non zero elements decreases)
- Error increases and Performance accuarcy drops(model starts underfitting)
- Run Time is also fast as sparsity increases

[8.3.3] Using RandomSearch CV :

L2 Regularization :

In [146]:

```
%%time  
if __name__ == "__main__":  
    LR_randomTrain(X_train_tfidfuni_std, y_train, penalty = 'l2')
```

Optimal C: {'C': 0.8189852652941816}

CrossValidation Error: 0.074


CrossValidation Accuracy: 92.598 %

CPU times: user 7min 51s, sys: 936 ms, total: 7min 52s

Wall time: 7min 52s


In [148]:

```
%%time
if __name__ == "__main__":
    LR_Test(X_train_tfidfuni_std, X_test_tfidfuni_std, y_train, y_test, 'l2', 0.818)
```



```
[[13140  5907]
 [ 2973 87060]]
```

Test Error : 0.081
Test Accuracy : 91.859 %
True Negative : 13140
False Positive : 5907
False Negative : 2973
True Positive : 87060
Precision Score : 0.876
Recall Score : 0.828
F1 Score : 0.849



CPU times: user 6.42 s, sys: 32 ms, total: 6.45 s
Wall time: 6.1 s

L1 Regularization :

In [147]:

```
%%time  
if __name__ == "__main__":  
    LR_randomTrain(X_train_tfidfuni_std, y_train, penalty = 'l1')
```

Optimal C: {'C': 0.01676258606542679}


CrossValidation Error: 0.073

CrossValidation Accuracy: 92.681 %

CPU times: user 9min 15s, sys: 2.16 s, total: 9min 17s
Wall time: 9min 17s


In [149]:

```
%%time
if __name__ == "__main__":
    LR_Test(X_train_tfidfuni_std, X_test_tfidfuni_std, y_train, y_test, 'l1', 0.016)
```



```
[[12930  6117]
 [ 2702 87331]]
```

Test Error : 0.081
Test Accuracy : 91.915 %
True Negative : 12930
False Positive : 6117
False Negative : 2702
True Positive : 87331
Precision Score : 0.881
Recall Score : 0.824
F1 Score : 0.849



CPU times: user 6.04 s, sys: 28 ms, total: 6.06 s
Wall time: 5.69 s

[8.3.4] MultiCollinearity :

In [150]:

```
if __name__ == "__main__":  
    multicollinear_check(X_train_tfidfuni_std, X_test_tfidfuni_std, y_train, y_test,0.  
001)
```

-----BEFORE PERTUBATION TEST-----

Sample Weights:	0	1	2	3	4	5	
6 \							
0	0.006177	0.054816	0.039201	0.00835	0.005969	-0.00787	0.009118
7		8	9	...	3868	3869	3870 \
0	-0.008688	0.008616	0.014231	...	0.014822	0.010729	0.005334
3871		3872	3873	3874	3875	3876	3877
0	-0.003729	0.006976	0.023702	0.023845	0.006898	0.018194	-0.002004

[1 rows x 3878 columns]

Size of weight vector: 3878

Non zero weights: 3878

Test Accuracy : 91.825 %

-----AFTER PERTUBATION TEST-----

Sample Weights:	0	1	2	3	4		
5							
6 \							
0	0.006186	0.054843	0.039186	0.008354	0.005992	-0.007852	0.00912
7		8	9	...	3868	3869	3870 \
0	-0.008659	0.008639	0.014248	...	0.014836	0.010734	0.005318
3871		3872	3873	3874	3875	3876	3877
0	-0.003714	0.006912	0.023699	0.023831	0.006924	0.01818	-0.002012

[1 rows x 3878 columns]

Size of weight vector: 3878

Non zero weights: 3878

Test Accuracy : 91.828 %

Number of features with weights changing greater than 30% : 3

Following are the 3 features that are multicollinear

3154 3698 3828

[8.3.5] FeatureImportance :

In [149]:

```
if __name__ == "__main__":  
    feature_importance(X_train_tfidfuni_std,X_test_tfidfuni_std,y_train,y_test,0.001,tfidf_unigram)
```

-----Top 25 Negative Words with high Importance-----

Coefficient	Factor	Features
-0.347543		not
-0.242920		disappoint
-0.209474		worst
-0.162968		aw
-0.156885		terribl
-0.145698		horribl
-0.143453		return
-0.134956		tast
-0.126144		threw
-0.119785		money
-0.119123		unfortun
-0.116417		stale
-0.112915		didnt
-0.111654		wast
-0.110265		bland
-0.108177		thought
-0.106271		disgust
-0.103017		even
-0.099548		would
-0.098077		hope
-0.097126		weak
-0.094650		mayb
-0.094515		noth
-0.093463		yuck
-0.091183		tasteless

-----Top 25 Positive Words with high Importance-----

Coefficient	Factor	Features
0.156612		fantast
0.160048		beat
0.160520		satisfi
0.162330		glad
0.167817		enjoy
0.176734		yummi
0.178656		smooth
0.180394		addict
0.182064		awesom
0.191334		thank
0.192106		happi
0.195537		find
0.201394		wonder
0.201583		easi
0.210871		tasti
0.217372		amaz
0.245529		favorit
0.249742		nice
0.310194		excel
0.345013		perfect
0.385975		good
0.411531		delici
0.459631		best
0.489395		love
0.660038		great

[8.4] TF-IDF(bigram) :

In [11]:

```
%%time
tfidf_bigram = TfidfVectorizer(ngram_range=(1, 2),min_df = 0.0005)
X_train_tfidfbi = tfidf_bigram.fit_transform(X_train)
print("Type of Count Vectorizer: ",type(X_train_tfidfbi))
print("The shape of text TFIDF vectorizer: ", X_train_tfidfbi.get_shape())
print("Number of unique word: ", X_train_tfidfbi.get_shape()[1])
```

```
Type of Count Vectorizer: <class 'scipy.sparse.csr.csr_matrix'>
The shape of text TFIDF vectorizer: (254519, 10711)
Number of unique word: 10711
CPU times: user 36.7 s, sys: 428 ms, total: 37.1 s
Wall time: 37.1 s
```

In [30]:

```
%%time
X_test_tfidfbi = tfidf_bigram.transform(X_test)
print("The shape of text TFIDF vectorizer: ", X_test_tfidfbi.get_shape())
print("Number of unique word: ", X_test_tfidfbi.get_shape()[1])
```

```
The shape of text TFIDF vectorizer: (109080, 10711)
Number of unique word: 10711
CPU times: user 11.2 s, sys: 28 ms, total: 11.2 s
Wall time: 11.2 s
```

In [152]:

```
dumpfile(X_train_tfidfbi,"X_train_tfidfbi")
dumpfile(X_test_tfidfbi,"X_test_tfidfbi")
```

In [151]:

```
X_train_tfidfbi = loadfile("X_train_tfidfbi")
X_test_tfidfbi = loadfile("X_test_tfidfbi")
```

In [152]:

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler(with_mean=False)
X_train_tfidfbi_std = sc.fit_transform(X_train_tfidfbi)
```

In [153]:

```
X_test_tfidfbi_std = sc.transform(X_test_tfidfbi)
```

In [154]:

```
print("Shape of Training Data: ",X_train_tfidfbi_std.shape)
print("Shape of Test Data: ",X_test_tfidfbi_std.shape)
```

```
Shape of Training Data: (254519, 10711)
Shape of Test Data: (109080, 10711)
```

[8.4.1] Using GridSearch CV :

L2 Regularization :


In [157]:

```
%%time
if __name__ == "__main__":
    LR_gridTrain(X_train_tfidfbi_std, y_train, penalty = 'l2')
```

Optimal C: {'C': 0.001}

CrossValidation Error: 0.063

CrossValidation Accuracy: 93.67 %



CV Error for each value of C: [0.065 0.063 0.066 0.066 0.067 0.067 0.067
0.066 0.067 0.067]


CV Accuracy for each value of C: [0.935 0.937 0.934 0.934 0.933 0.933 0.9
33 0.934 0.933 0.933]

CPU times: user 1h 7min 34s, sys: 30 s, total: 1h 8min 4s

Wall time: 31min 42s


In [171]:

```
%%time
if __name__ == "__main__":
    LR_Test(X_train_tfidfbi_std, X_test_tfidfbi_std, y_train, y_test, 'l2', 0.001)
```



```
[[14255  4792]
 [ 2663 87370]]
```

Test Error : 0.068
 Test Accuracy : 93.166 %
 True Negative : 14255
 False Positive : 4792
 False Negative : 2663
 True Positive : 87370
 Precision Score : 0.895
 Recall Score : 0.859
 F1 Score : 0.876



CPU times: user 11.6 s, sys: 104 ms, total: 11.7 s
 Wall time: 6.09 s

L1 Regularization :


In [158]:

```
%%time
if __name__ == "__main__":
    LR_gridTrain(X_train_tfidfbi_std, y_train, penalty = 'l1')
```

Optimal C: {'C': 0.01}

CrossValidation Error: 0.061

CrossValidation Accuracy: 93.89 %



CV Error for each value of C: [0.149 0.083 0.061 0.064 0.066 0.066 0.066 0.066 0.066 0.066]


CV Accuracy for each value of C: [0.851 0.917 0.939 0.936 0.934 0.934 0.934 0.934 0.934 0.934]

CPU times: user 11min 36s, sys: 5.46 s, total: 11min 42s

Wall time: 11min 41s


In [172]:

```
%%time
if __name__ == "__main__":
    LR_Test(X_train_tfidfbi_std, X_test_tfidfbi_std, y_train, y_test, 'l1', 0.01)
```



```
[[13968  5079]
 [ 2167 87866]]
```

Test Error : 0.066
Test Accuracy : 93.357 %
True Negative : 13968
False Positive : 5079
False Negative : 2167
True Positive : 87866
Precision Score : 0.906
Recall Score : 0.855
F1 Score : 0.877



CPU times: user 6.5 s, sys: 48 ms, total: 6.54 s
Wall time: 6.14 s

[8.4.2] Sparsity And Accuracy on Weight Vectors(L1 Regularization) :

In [33]:

```
if __name__ == "__main__":  
    sparsity_check(X_train_tfidfbi_std, X_test_tfidfbi_std ,y_train ,y_test)
```

Sparsity and Accuarcy when C = 10
Number of non-zero weighhts: 10698
Train Error: 0.045
Test Error: 0.07
Test Accuracy : 0.92995
Run Time :9.959680000000006 sec

Sparsity and Accuarcy when C = 1
Number of non-zero weighhts: 10608
Train Error: 0.045
Test Error: 0.07
Test Accuracy : 0.9301
Run Time :8.638589999999994 sec

Sparsity and Accuarcy when C = 0.1
Number of non-zero weighhts: 9956
Train Error: 0.046
Test Error: 0.069
Test Accuracy : 0.93129
Run Time :6.357533000000004 sec

Sparsity and Accuarcy when C = 0.01
Number of non-zero weighhts: 5396
Train Error: 0.052
Test Error: 0.066
Test Accuracy : 0.93356
Run Time :5.3310010000000005 sec

Sparsity and Accuarcy when C = 0.001
Number of non-zero weighhts: 539
Train Error: 0.08
Test Error: 0.089
Test Accuracy : 0.91087
Run Time :2.646493999999999 sec

Sparsity and Accuarcy when C = 0.0001
Number of non-zero weighhts: 14
Train Error: 0.149
Test Error: 0.174
Test Accuracy : 0.82603
Run Time :1.4733790000000084 sec

Observation : Here $C = 1/\lambda$, we can see as C decreases(λ increases)

- Sparsity Increases(Number of non zero elements decreases)
- Error increases and Performance accuarcy drops(model starts underfitting)
- Run Time is also fast as sparsity increases

[8.4.3] Using RandomSearch CV :

L2 Regularization :

In []:

```
%%time
if __name__ == "__main__":
    LR_randomTrain(X_train_tfidfbi_std, y_train, penalty = 'l2')
```

Optimal C: {'C': 0.09716463627118838}

CrossValidation Error: 0.066


CrossValidation Accuracy: 93.366 %

CPU times: user 1h 24min 43s, sys: 35.6 s, total: 1h 25min 18s

Wall time: 39min 50s


In [36]:

```
%%time
if __name__ == "__main__":
    LR_Test(X_train_tfidfbi_std, X_test_tfidfbi_std, y_train, y_test, 'l2',0.097)
```

```
[[14387 4660]  
 [ 2977 87056]]
```

Test Error : 0.07
Test Accuracy : 92.999 %
True Negative : 14387
False Positive : 4660
False Negative : 2977
True Positive : 87056
Precision Score : 0.889
Recall Score : 0.861
F1 Score : 0.874



CPU times: user 30.9 s, sys: 200 ms, total: 31.1 s
Wall time: 15.5 s

L1 Regularization :

In []:

```
%%time  
if __name__ == "__main__":  
    LR_randomTrain(X_train_tfidfbi_std, y_train, penalty = 'l1')
```

Optimal C: {'C': 0.01845898506776376}


CrossValidation Error: 0.061

CrossValidation Accuracy: 93.904 %

CPU times: user 14min 6s, sys: 4.92 s, total: 14min 10s
Wall time: 14min 11s


In [34]:

```
%%time
if __name__ == "__main__":
    LR_Test(X_train_tfidfbi_std, X_test_tfidfbi_std, y_train, y_test, 'l1', 0.018)
```



```
[[14190  4857]
 [ 2406 87627]]
```

Test Error : 0.067
 Test Accuracy : 93.342 %
 True Negative : 14190
 False Positive : 4857
 False Negative : 2406
 True Positive : 87627
 Precision Score : 0.901
 Recall Score : 0.859
 F1 Score : 0.878



CPU times: user 7.26 s, sys: 44 ms, total: 7.31 s
 Wall time: 7.02 s

[8.4.4] MultiCollinearity :

In [155]:

```
if __name__ == "__main__":
    multicollinear_check(X_train_tfidfbi_std, X_test_tfidfbi_std, y_train, y_test,0.00
1)
```

-----BEFORE PERTUBATION TEST-----

```
Sample Weights:      0      1      2      3      4      5
      6      \
0 -0.000201  0.039059  0.03512 -0.02023 -0.002252  0.035492  0.009385

      7      8      9      ...      10701      10702      10703  \
0 -0.01361  0.039944  0.014182  ...      -0.000493  0.011343  0.026953

      10704      10705      10706      10707      10708      10709      10710
0  0.030759  0.008869 -0.010293  0.024731  0.012831  0.008689 -0.00666
```

```
[1 rows x 10711 columns]
Size of weight vector: 10711
Non zero weights: 10711
Test Accuracy : 93.166 %
```

-----AFTER PERTUBATION TEST-----

```
Sample Weights:      0      1      2      3      4      5
      6      \
0 -0.000221  0.039054  0.035148 -0.020233 -0.002193  0.035475  0.009385

      7      8      9      ...      10701      10702      10703  \
0 -0.013613  0.039935  0.014227  ...      -0.000511  0.01135  0.026948

      10704      10705      10706      10707      10708      10709      10710
0  0.030755  0.008834 -0.010282  0.024704  0.012802  0.008696 -0.006638
```

```
[1 rows x 10711 columns]
Size of weight vector: 10711
Non zero weights: 10711
Test Accuracy : 93.168 %
```

Number of features with weights changing greater than 30% : 33

Following are the 33 features that are multicollinear

```
379 759 781 991 1036 1572 1742 2069 2984 3017 3584 3974 3989 4728 4898 527
2 6226 6484 6650 7382 7920 8095 8323 8532 8645 8819 9052 9204 9255 9256 95
14 9560 10071
```

[8.4.5] FeatureImportance :

In [156]:

```
if __name__ == "__main__":  
    feature_importance(X_train_tfidfbi_std,X_test_tfidfbi_std,y_train,y_test,0.001,tfidf_bigram)
```

-----Top 25 Negative Words with high Importance-----

Coefficient	Factor	Features
-0.309475		disappoint
-0.190949		not
-0.179745		worst
-0.168787		not worth
-0.154740		not good
-0.147985		aw
-0.147965	not	recommend
-0.147495		two star
-0.147186		terribl
-0.137424		not buy
-0.135074		horribl
-0.120756		stale
-0.118112		return
-0.116322		threw
-0.116124		unfortun
-0.112878		weak
-0.111305		wont buy
-0.111163		bad
-0.110713		bland
-0.106437	wast	money
-0.104984		disgust
-0.092486		mayb
-0.091960		sorri
-0.090592		tast
-0.090431		tasteless

-----Top 25 Positive Words with high Importance-----

Coefficient	Factor	Features
0.133896		glad
0.138518		fantast
0.139108		hook
0.140077		enjoy
0.143920		easi
0.156743		yummi
0.163326		wonder
0.163776		thank
0.164907		addict
0.165519		awesom
0.167443		satisfi
0.181310		amaz
0.194191		nice
0.196270		favorit
0.211577	not	disappoint
0.216029	high	recommend
0.217148		tasti
0.227297		happi
0.243336		excel
0.264744		perfect
0.344418		best
0.354443		delici
0.360684		good
0.382252		love
0.453487		great

[8.5] Average Word2Vec :

In [13]:

```
i=0
list_of_sent_train=[]
for sent in X_train:
    list_of_sent_train.append(sent.split())
```

In [14]:

```
print(X_train[5])
print("*****")
print(list_of_sent_train[5])
```

bought apart infest fruit fli hour trap mani fli within day practic gone m
ay not long term solut fli drive crazi consid buy one surfac sticki tri av
oid touch

**

```
['bought', 'apart', 'infest', 'fruit', 'fli', 'hour', 'trap', 'mani', 'fl  
i', 'within', 'day', 'practic', 'gone', 'may', 'not', 'long', 'term', 'sol  
ut', 'fli', 'drive', 'crazi', 'consid', 'buy', 'one', 'surfac', 'sticki',  
'tri', 'avoid', 'touch']
```

In [15]:

```
%%time
## Word2Vec Model considering only those words that occur atleast 5 times in the corpus
min_count = 5
w2v_model = Word2Vec(list_of_sent_train, min_count = min_count, size = 50, workers = 4)
w2v_words = list(w2v_model.wv.vocab)
```

CPU times: user 2min 8s, sys: 240 ms, total: 2min 9s

Wall time: 36 s

In [16]:

```
i=0
list_of_sent_test=[]
for sent in X_test:
    list_of_sent_test.append(sent.split())
```

In [16]:

```
print(X_test[5])
print("*****")
print(list_of_sent_test[5])
```

dog love love dog food say love twice dog food one dog favorit

**

```
['dog', 'love', 'love', 'dog', 'food', 'say', 'love', 'twice', 'dog', 'foo  
d', 'one', 'dog', 'favorit']
```


In [37]:

```
%%time
X_train_avgw2v = [] # the avg-w2v for each sentence/review is stored in this list
for sent in list_of_sent_train:
    sent_vec = np.zeros(50) # as word vectors are of zero length
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    X_train_avgw2v.append(sent_vec)
```

CPU times: user 7min 47s, sys: 64 ms, total: 7min 47s
Wall time: 7min 47s

In [38]:

```
%%time
X_test_avgw2v = [] # the avg-w2v for each sentence/review is stored in this list
for sent in list_of_sent_test:
    sent_vec = np.zeros(50) # as word vectors are of zero length
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    X_test_avgw2v.append(sent_vec)
```

CPU times: user 3min 33s, sys: 12 ms, total: 3min 33s
Wall time: 3min 33s

In [68]:

```
#Checking NAN in test data if any
np.any(np.isnan(X_test_avgw2v))
```

Out[68]:

False

In [41]:

```
dumpfile(X_train_avgw2v,"X_train_avgw2v")
dumpfile(X_test_avgw2v,"X_test_avgw2v")
```

In [157]:

```
X_train_avgw2v = loadfile("X_train_avgw2v")
X_test_avgw2v = loadfile("X_test_avgw2v")
```

In [158]:

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler(with_mean=False)
X_train_avgw2v_std = sc.fit_transform(X_train_avgw2v)
```

In [159]:

```
X_test_avgw2v_std = sc.transform(X_test_avgw2v)
```

In [160]:

```
print("Number of rows in Train Data: ",len(X_train_avgw2v_std))
print("Number of features in Train Data: ",len(X_train_avgw2v_std[0]))
print("Number of rows in Test Data: ",len(X_test_avgw2v_std))
print("Number of features in Test Data: ",len(X_test_avgw2v_std[0]))
```

Number of rows in Train Data: 254519

Number of features in Train Data: 50

Number of rows in Test Data: 109080

Number of features in Test Data: 50

[8.5.1] Using GridSearch CV :

L2 Regularization :


In [49]:

```
%%time
if __name__ == "__main__":
    LR_gridTrain(X_train_avgw2v_std, y_train, penalty = 'l2')
```

Optimal C: {'C': 50}

CrossValidation Error: 0.097

CrossValidation Accuracy: 90.291 %



CV Error for each value of C: [0.108 0.098 0.097 0.097 0.097 0.097 0.097 0.097 0.097 0.097]


CV Accuracy for each value of C: [0.892 0.902 0.903 0.903 0.903 0.903 0.903 0.903 0.903 0.903]

CPU times: user 8min 2s, sys: 1.04 s, total: 8min 3s

Wall time: 8min


In [50]:

```
%%time
if __name__ == "__main__":
    LR_Test(X_train_avg2v_std, X_test_avg2v_std, y_train, y_test, 'l2', 50)
```



```
[[10342  8705]
 [ 3034 86999]]
```

Test Error : 0.108
Test Accuracy : 89.238 %
True Negative : 10342
False Positive : 8705
False Negative : 3034
True Positive : 86999
Precision Score : 0.841
Recall Score : 0.755
F1 Score : 0.787



CPU times: user 7.46 s, sys: 12 ms, total: 7.47 s
Wall time: 7.03 s

L1 Regularization :


In [51]:

```
%%time
if __name__ == "__main__":
    LR_gridTrain(X_train_avgw2v_std, y_train, penalty = 'l1')
```

Optimal C: {'C': 1000}

CrossValidation Error: 0.097

CrossValidation Accuracy: 90.29 %



CV Error for each value of C: [0.135 0.102 0.097 0.097 0.097 0.097 0.097 0.097 0.097 0.097]


CV Accuracy for each value of C: [0.865 0.898 0.903 0.903 0.903 0.903 0.903 0.903 0.903 0.903]

CPU times: user 1h 35min 19s, sys: 3.26 s, total: 1h 35min 22s

Wall time: 1h 35min 19s


In [53]:

```
%%time
if __name__ == "__main__":
    LR_Test(X_train_avg2v_std, X_test_avg2v_std, y_train, y_test, 'l1',1000)
```



```
[[10345  8702]  
 [ 3027 87006]]
```

Test Error : 0.108
Test Accuracy : 89.247 %
True Negative : 10345
False Positive : 8702
False Negative : 3027
True Positive : 87006
Precision Score : 0.841
Recall Score : 0.755
F1 Score : 0.788



CPU times: user 2min 18s, sys: 44 ms, total: 2min 18s
Wall time: 2min 18s

[8.5.2] Sparsity And Accuracy on Weight Vectors(L1 Regularization) :

In [52]:

```
if __name__ == "__main__":  
    sparsity_check(X_train_avg2v_std, X_test_avg2v_std ,y_train ,y_test)
```

Sparsity and Accuracy when C = 10
Number of non-zero weights: 50
Train Error: 0.097
Test Error: 0.107
Test Accuracy : 0.89251
Run Time :77.11510200000066 sec

Sparsity and Accuracy when C = 1
Number of non-zero weights: 50
Train Error: 0.097
Test Error: 0.107
Test Accuracy : 0.8925
Run Time :76.87039000000004 sec

Sparsity and Accuracy when C = 0.1
Number of non-zero weights: 50
Train Error: 0.097
Test Error: 0.108
Test Accuracy : 0.89241
Run Time :93.50651099999959 sec

Sparsity and Accuracy when C = 0.01
Number of non-zero weights: 49
Train Error: 0.097
Test Error: 0.108
Test Accuracy : 0.89203
Run Time :30.468505000000732 sec

Sparsity and Accuracy when C = 0.001
Number of non-zero weights: 38
Train Error: 0.101
Test Error: 0.113
Test Accuracy : 0.88731
Run Time :4.664448999999877 sec

Sparsity and Accuracy when C = 0.0001
Number of non-zero weights: 13
Train Error: 0.133
Test Error: 0.151
Test Accuracy : 0.84876
Run Time :1.6230040000009467 sec

Observation : Here $C = 1/\lambda$, we can see as C decreases(λ increases)

- Sparsity Increases(Number of non zero elements decreases)
- Error increases and Performance accuracy drops(model starts underfitting)
- Run Time is also fast as sparsity increases

[8.5.3] Using RandomSearch CV :

L2 Regularization :

In [54]:

```
%%time
if __name__ == "__main__":
    LR_randomTrain(X_train_avgw2v_std, y_train, penalty = 'l2')
```

Optimal C: {'C': 2.6804844914828405}

CrossValidation Error: 0.097


CrossValidation Accuracy: 90.293 %

CPU times: user 8min 47s, sys: 1.09 s, total: 8min 48s

Wall time: 8min 45s


In [72]:

```
%%time
if __name__ == "__main__":
    LR_Test(X_train_avg2v_std, X_test_avg2v_std, y_train, y_test, 'l2', 2.680)
```



```
[[10345  8702]
 [ 3033 87000]]
```

Test Error : 0.108
 Test Accuracy : 89.242 %
 True Negative : 10345
 False Positive : 8702
 False Negative : 3033
 True Positive : 87000
 Precision Score : 0.841
 Recall Score : 0.755
 F1 Score : 0.787



CPU times: user 10 s, sys: 4 ms, total: 10 s
 Wall time: 9.58 s

L1 Regularization :

In []:

```
%%time
if __name__ == "__main__":
    LR_randomTrain(X_train_avgw2v_std, y_train, penalty = 'l1')
```

Optimal C: {'C': 0.17690344265644772}

CrossValidation Error: 0.097


CrossValidation Accuracy: 90.288 %

CPU times: user 1h 53min 48s, sys: 5.83 s, total: 1h 53min 54s

Wall time: 1h 53min 51s


In [15]:

```
%%time
if __name__ == "__main__":
    LR_Test(X_train_avg2v_std, X_test_avg2v_std, y_train, y_test, 'l1', 0.176)
```



```
[[10327  8720]
 [ 3011 87022]]
```

Test Error : 0.108
 Test Accuracy : 89.246 %
 True Negative : 10327
 False Positive : 8720
 False Negative : 3011
 True Positive : 87022
 Precision Score : 0.842
 Recall Score : 0.754
 F1 Score : 0.787



CPU times: user 1min 23s, sys: 148 ms, total: 1min 23s
Wall time: 1min 23s

[8.5.4] MultiCollinearity :

In [161]:

```
if __name__ == "__main__":  
    multicollinear_check(X_train_avg2v_std,X_test_avg2v_std,y_train,y_test,50)
```

-----BEFORE PERTUBATION TEST-----

```
Sample Weights:          0          1          2          3          4  
5          6  \  
0 -0.588587  0.446414 -0.533977 -0.788838 -0.422129 -0.363002  0.372933  
  
          7          8          9      ...          40          41          42  \  
0  0.04856  0.144315 -0.020016      ...      -0.237664 -0.42991 -0.534093  
  
          43          44          45          46          47          48          49  
0 -0.174668 -0.363496  0.07075  0.365204 -0.169001 -0.25998 -0.458187
```

[1 rows x 50 columns]
Size of weight vector: 50
Non zero weights: 50
Test Accuracy : 89.238 %

-----AFTER PERTUBATION TEST-----

```
Sample Weights:          0          1          2          3          4  
5          6  \  
0 -0.576476  0.44317 -0.519533 -0.756785 -0.42583 -0.356126  0.35971  
  
          7          8          9      ...          40          41          42  \  
0  0.046812  0.136165 -0.012172      ...      -0.227333 -0.428892 -0.525311  
  
          43          44          45          46          47          48          49  
0 -0.188625 -0.345497  0.076696  0.360842 -0.160229 -0.258166 -0.440162
```

[1 rows x 50 columns]
Size of weight vector: 50
Non zero weights: 50
Test Accuracy : 89.248 %

Number of features with weights changing greater than 30% : 3

Following are the 3 features that are multicollinear
9 17 26

[8.6] TF-IDF Weighted Word2Vec :

In [17]:

```
%%time
tfidf_feat = tfidf_bigram.get_feature_names() # tfidf words/col-names
X_train_tfidf2v = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in list_of_sent_train:
    sent_vec = np.zeros(50)
    weight_sum =0;
    for word in sent:
        try:
            vec = w2v_model.wv[word]
            tfidf = X_train_tfidfbi[row, tfidf_feat.index(word)]
            sent_vec += (vec * tfidf)
            weight_sum += tfidf
        except:
            pass

    if weight_sum != 0:
        sent_vec /= weight_sum
    X_train_tfidf2v.append(sent_vec)
    row += 1
```

CPU times: user 29min 29s, sys: 876 ms, total: 29min 30s

Wall time: 29min 29s

In [37]:

```
%%time
X_test_tfidf2v = [];
row=0;
for sent in list_of_sent_test:
    sent_vec = np.zeros(50)
    weight_sum =0;
    for word in sent:
        try:
            vec = w2v_model.wv[word]
            tfidf = X_train_tfidfbi[row, tfidf_feat.index(word)]
            sent_vec += (vec * tfidf)
            weight_sum += tfidf
        except:
            pass

    if weight_sum != 0:
        sent_vec /= weight_sum
    X_test_tfidf2v.append(sent_vec)
    row += 1
```

CPU times: user 12min 57s, sys: 2.1 s, total: 13min

Wall time: 12min 59s

In [38]:

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler(with_mean=False)
X_train_tfidf2v_std = sc.fit_transform(X_train_tfidf2v)
```

In [39]:

```
X_test_tfidfw2v_std = sc.transform(X_test_tfidfw2v)
```

In [40]:

```
print("Number of rows in Train Data: ",len(X_train_tfidfw2v_std))
print("Number of features in Train Data: ",len(X_train_tfidfw2v_std[0]))
print("Number of rows in Test Data: ",len(X_test_tfidfw2v_std))
print("Number of features in Test Data: ",len(X_test_tfidfw2v_std[0]))
```

```
Number of rows in Train Data: 254519
Number of features in Train Data: 50
Number of rows in Test Data: 109080
Number of features in Test Data: 50
```

In [41]:

```
dumpfile(X_train_tfidfw2v_std,"X_train_tfidfw2v_std")
dumpfile(X_test_tfidfw2v_std,"X_test_tfidfw2v_std")
```

[8.6.1] Using GridSearch CV :

L2 Regularization :


In [42]:

```
%%time
if __name__ == "__main__":
    LR_gridTrain(X_train_tfidfw2v_std, y_train, penalty = 'l2')
```

Optimal C: {'C': 0.0001}

CrossValidation Error: 0.149

CrossValidation Accuracy: 85.075 %



CV Error for each value of C: [0.149 0.149 0.149 0.149 0.149 0.149 0.149 0.149 0.149 0.149]


CV Accuracy for each value of C: [0.851 0.851 0.851 0.851 0.851 0.851 0.851 0.851 0.851 0.851]

CPU times: user 32.8 s, sys: 1.21 s, total: 34 s

Wall time: 30.9 s


In [43]:

```
%%time
if __name__ == "__main__":
    LR_Test(X_train_tfidf2v_std, X_test_tfidf2v_std, y_train, y_test, 'l2', 0.0001)
```



```
[[ 0 19047]
 [ 0 90033]]
```

Test Error : 0.175
Test Accuracy : 82.539 %
True Negative : 0
False Positive : 19047
False Negative : 0
True Positive : 90033
Precission Score : 0.413
Recall Score : 0.5
F1 Score : 0.452



CPU times: user 1.41 s, sys: 8 ms, total: 1.42 s
Wall time: 990 ms

L1 Regularization :


In [47]:

```
%%time
if __name__ == "__main__":
    LR_gridTrain(X_train_tfidfw2v_std, y_train, penalty = 'l1')
```

Optimal C: {'C': 0.0001}

CrossValidation Error: 0.149

CrossValidation Accuracy: 85.075 %



CV Error for each value of C: [0.149 0.149 0.149 0.149 0.149 0.149 0.149 0.149 0.149 0.149]


CV Accuracy for each value of C: [0.851 0.851 0.851 0.851 0.851 0.851 0.851 0.851 0.851 0.851]

CPU times: user 28.1 s, sys: 1.07 s, total: 29.2 s

Wall time: 25.7 s


In [48]:

```
%%time
if __name__ == "__main__":
    LR_Test(X_train_tfidf2v_std, X_test_tfidf2v_std, y_train, y_test, 'l1', 0.0001)
```



```
[[ 0 19047]
 [ 0 90033]]
```

Test Error : 0.175
Test Accuracy : 82.539 %
True Negative : 0
False Positive : 19047
False Negative : 0
True Positive : 90033
Precision Score : 0.413
Recall Score : 0.5
F1 Score : 0.452



CPU times: user 1.34 s, sys: 28 ms, total: 1.36 s
Wall time: 962 ms

[8.6.2] Sparsity And Accuracy on Weight Vectors(L1 Regularization) :

In [44]:

```
if __name__ == "__main__":  
    sparsity_check(X_train_tfidfv2v_std, X_test_tfidfv2v_std ,y_train ,y_test)
```

Sparsity and Accuracy when C = 10
Number of non-zero weights: 0
Train Error: 0.149
Test Error: 0.175
Test Accuracy : 0.82539
Run Time :0.4394509999992806 sec

Sparsity and Accuracy when C = 1
Number of non-zero weights: 0
Train Error: 0.149
Test Error: 0.175
Test Accuracy : 0.82539
Run Time :0.4273480000001655 sec

Sparsity and Accuracy when C = 0.1
Number of non-zero weights: 0
Train Error: 0.149
Test Error: 0.175
Test Accuracy : 0.82539
Run Time :0.4278549999999086 sec

Sparsity and Accuracy when C = 0.01
Number of non-zero weights: 0
Train Error: 0.149
Test Error: 0.175
Test Accuracy : 0.82539
Run Time :0.4308789999995497 sec

Sparsity and Accuracy when C = 0.001
Number of non-zero weights: 0
Train Error: 0.149
Test Error: 0.175
Test Accuracy : 0.82539
Run Time :0.4345000000002983 sec

Sparsity and Accuracy when C = 0.0001
Number of non-zero weights: 0
Train Error: 0.149
Test Error: 0.175
Test Accuracy : 0.82539
Run Time :0.4260019999992437 sec

[8.6.3] Using RandomSearch CV :

L2 Regularization :

In [45]:

```
%%time
if __name__ == "__main__":
    LR_randomTrain(X_train_tfidfv2v_std, y_train, penalty = 'l2')
```

Optimal C: {'C': 1.2928165589093843}

CrossValidation Error: 0.149


CrossValidation Accuracy: 85.075 %

CPU times: user 32.4 s, sys: 980 ms, total: 33.4 s

Wall time: 30.2 s


In [46]:

```
%%time
if __name__ == "__main__":
    LR_Test(X_train_tfidf2v_std, X_test_tfidf2v_std, y_train, y_test, 'l2', 1.29)
```



```
[[ 0 19047]
 [ 0 90033]]
```

Test Error : 0.175
Test Accuracy : 82.539 %
True Negative : 0
False Positive : 19047
False Negative : 0
True Positive : 90033
Precission Score : 0.413
Recall Score : 0.5
F1 Score : 0.452



CPU times: user 1.36 s, sys: 8 ms, total: 1.36 s
Wall time: 981 ms

L1 Regularization :

In [49]:

```
%%time  
if __name__ == "__main__":  
    LR_randomTrain(X_train_tfidf2v_std, y_train, penalty = 'l1')
```

Optimal C: {'C': 0.9964195542921448}


CrossValidation Error: 0.149

CrossValidation Accuracy: 85.075 %

CPU times: user 27.4 s, sys: 1.16 s, total: 28.5 s
Wall time: 25.4 s


In [51]:

```
%%time
if __name__ == "__main__":
    LR_Test(X_train_tfidf2v_std, X_test_tfidf2v_std, y_train, y_test, 'l1', 0.99)
```



```
[[ 0 19047]
 [ 0 90033]]
```

Test Error : 0.175
Test Accuracy : 82.539 %
True Negative : 0
False Positive : 19047
False Negative : 0
True Positive : 90033
Precision Score : 0.413
Recall Score : 0.5
F1 Score : 0.452



CPU times: user 1.34 s, sys: 28 ms, total: 1.37 s
Wall time: 954 ms

Observation : Tfidf Weighted Word2Vec classifies all the test data(unseen points) as Positive(majority class). It is like a dumb model as we can see from the ROC plot,it overlaps with simple model and True Negative is also zero.

[9] Conclusion :

Grid Search Cross Validation :

Featurization Model	L2 Regularization				L1 Regularization			
	Accuracy	Precission	Recall	F1 score	Accuracy	Precission	Recall	F1 score
BOW(unigram)	91.569 %	0.878	0.813	0.84	91.548 %	0.875	0.816	0.841
BOW(bigram)	93.004 %	0.908	0.838	0.868	93.039 %	0.905	0.843	0.87
TF-IDF(unigram)	91.823 %	0.878	0.824	0.847	91.892 %	0.877	0.828	0.85
TF-IDF(bigram)	93.166 %	0.895	0.859	0.876	93.357 %	0.906	0.855	0.877
Average Word2vec	89.238 %	0.841	0.755	0.787	89.247 %	0.841	0.755	0.788
TF-IDF Wweighted Word2Vec	82.539 %	0.413	0.5	0.452	82.539 %	0.413	0.5	0.452

Random Search Cross Validation :

Featurization Model	L2 Regularization				L1 Regularization			
	Accuracy	Precission	Recall	F1 score	Accuracy	Precission	Recall	F1 score
BOW(unigram)	91.566 %	0.874	0.817	0.842	91.563 %	0.875	0.817	0.842
BOW(bigram)	92.794 %	0.886	0.856	0.87	92.92 %	0.89	0.856	0.872
TF-IDF(unigram)	91.859 %	0.876	0.828	0.849	91.915 %	0.881	0.824	0.849
TF-IDF(bigram)	92.999 %	0.889	0.861	0.874	93.342 %	0.901	0.859	0.878
Average Word2vec	89.242 %	0.841	0.755	0.787	89.246 %	0.842	0.754	0.787
TF-IDF Wweighted Word2Vec	82.539 %	0.413	0.5	0.452	82.539 %	0.413	0.5	0.452

1 - Tfidf with bigram performed the best with all performance metrics having an accuracy of 93.166 % and F1 Score of .876 in L2 regularization, accuracy of 93.357 % and F1 Score of .877 in L1 regularization.

2 - Both Gridserach and Random Search Crossvalidation gave almost equal results.

- 3 - L1 regularization creates a sparse weight vector ie all the less important features becomes zero.
- 4 - It is also observed that in L1 regularization, as hyperparameter C decreases(lambda increases),
- (i) Sparsity Increases(Number of non zero elements decreases)
 - (ii)Error increases and Performance accuracy drops(model starts underfitting)
 - (iii)Run Time is also fast as sparsity increases
- 5 - Logistic Regression gave better accuracies and results as compared to Naive Bayes model.
- 6 - Run Time Complexity of Logistic Regression is less. This model can be used for low latency applications.
- 7 - TF-IDF Wweighted Word2Vec performed worst in Logistic Regression model. It is like a dumb model as every unseen points gets classified as majority class(Positive).