

# Table of Contents

- 1 - Problem Statement
- 2 - Overview of Dataset
- 3 - Loading the Data
- 4 - Exploratory Data Analysis
  - 4.1 - Data Cleaning: Deduplication
- 5 - Text Preprocessing Using NLTK
- 6 - Train and Test Split of Data
- 7 - Naive Bayes Classification Model
  - 7.1 - Function to find the optimal alpha and error using 10-fold cross-validation
  - 7.2 - Function to find the features importance and prediction on TestData
- 8 - Featurization Methods
  - 8.1 - Bag Of Words(unigram)
  - 8.2 - Bag Of Words(bigram)
  - 8.3 - TF-IDF(unigram)
  - 8.4 - TF-IDF(bigram)
- 9 - Conclusion

## [1] Problem Statement :

- Time Based slicing to split Train Data(70%) and Test Data(30%).
- Applying Naive Bayes model to find the optimal alpha and using 10 fold Cross Validation in :
  - 1)Bag Of Words
  - 2)TF-IDF
- Finding the features Importance for each class label.
- Plotting Confusion Matrix.
- Using Precision, Recall, F1 Score as performance metrics and comparing between various featurization techniques.

## [2] Overview of Dataset :

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>  
(<https://www.kaggle.com/snap/amazon-fine-food-reviews>)

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454

Number of users: 256,059

Number of products: 74,258

Timespan: Oct 1999 - Oct 2012

Number of Attributes/Columns in data: 10

Attribute Information:

- 1.Id
- 2.ProductId - unique identifier for the product
- 3.UserId - unique identifier for the user
- 4.ProfileName
- 5.HelpfulnessNumerator - number of users who found the review helpful
- 6.HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
- 7.Score - rating between 1 and 5
- 8.Time - timestamp for the review
- 9.Summary - brief summary of the review
- 10.Text - text of the review

Objective: Given a review, determine whether the review is positive (Rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative?

[Ans] We could use the Score/Rating. A rating of 4 or 5 could be considered a positive review. A review of 1 or 2 could be considered negative. A review of 3 is neutral and ignored. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

## [3] Loading the Data :

In order to load the data, we have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

In [1]:

```
#Importing the necessary Packages
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import time
import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from IPython.display import HTML
from collections import OrderedDict

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle
```

In [2]:

```
import pickle

#dumping an object to file object using dump method
def dumpfile(a,file_Name):
    fileObject = open(file_Name,"wb")
    pickle.dump(a,fileObject,protocol=2)
    fileObject.close()

#Loading an object from file object using load method
def loadfile(file_Name):
    fileObject = open(file_Name,"rb")
    b = pickle.load(fileObject)
    return b
```

In [3]:

```
%%HTML
<style type="text/css">
table.dataframe td, table.dataframe th {
    border: 2px  black solid !important;
}
</style>
```

In [4]:

```
# using the SQLite Table to read data.
con = sqlite3.connect('database.sqlite')

#filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 """, con)
```

In [5]:

```
# Give reviews with Score>3 a positive rating, and reviews with a score<3 a negative rating.
def partition(x):
    if x < 3:
        return 'negative'
    return 'positive'

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
```

In [6]:

```
print("Number of datapoints: ",filtered_data.shape[0])
print("Number of attributes/features: ",filtered_data.shape[1])
HTML(filtered_data.head().to_html(index=False))
```

Number of datapoints: 525814

Number of attributes/features: 10

Out[6]:

Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	Helpfulr
1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1
2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0
3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	1
4	B000UA0QIQ	A395BORC6FGVXV	Karl	3	3
5	B006K2ZZ7K	A1UQRSCLF8GW1T	Michael D. Bigham "M. Wassir"	0	0



## [4] Exploratory Data Analysis :

## [4.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

**Deduplication 1:-** As can be seen below the same user has multiple reviews of the with the same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that

ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8)

ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8)

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

In [7]:

```
display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", con)

HTML(display.head().to_html(index=False))
```

Out[7]:

Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator
78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2	2
138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2	2
138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2	2
73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2	2
155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2	2



In [8]:

```
#Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False, kind='quicksort', na_position='last')
```

In [9]:

```
#Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep=
'first', inplace=False)
final.shape
```

Out[9]:

(364173, 10)

**Deduplication 2:-** It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

In [10]:

```
display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)

HTML(display.head().to_html(index=False))
```

Out[10]:

Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator
64422	B000MIDROQ	A161DK06JJMCYF	J. E. Stephens "Jeanne"	3	1
44737	B001EQ55RW	A2V0I904FH7ABY	Ram	3	2

In [11]:

```
final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
print(final.shape)
```

(364171, 10)



**Deduplication 3:-** It was also seen that a same user has given different reviews for a same product at same time. I think it is normal for a user to give multiple reviews about a product, but that should be in different time. So, all those rows with same user giving multiple reviews for a same product at same time are considered as duplicate and hence dropped.

In [12]:

```
display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId= "A8891HVRDJAM6"
ORDER BY ProductID
""", con)

HTML(display.head().to_html(index=False))
```

Out[12]:

Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator
86221	B000084E6V	A8891HVRDJAM6	Marfaux "Marfaux"	33	33
86236	B000084E6V	A8891HVRDJAM6	Marfaux "Marfaux"	3	3

In [13]:

```
final=final.drop_duplicates(subset={"ProductId","UserId","ProfileName","Time"}, keep='first', inplace=False)
print(final.shape)
```

(363633, 10)

**Deduplication 4:-** It was also seen that in few rows with Ids from 150493 to 150529 contain reviews regarding books, not fine foods. So I think these should be also removed from the dataset. After looking at the productid column, it can be noticed that all the observations for fine foods start with B followed by numbers except for Ids from 150493 to 150529. I suppose the reviews for book 'Chicken soup for the soul' have gotten into the dataset mistakenly as they contain the words "chicken soup."

In [14]:

```
display = final[final.ProductId == "0006641040"]  
HTML(display.head().to_html(index=False))
```

Out[14]:

Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator
150524	0006641040	ACITT7DI6IDDL	shari zychinski	0	0
150506	0006641040	A2IW4PEEKO2R0U	Tracy	1	1
150507	0006641040	A1S4A3IQ2MU7V4	sally sue "sally sue"	1	1
150508	0006641040	AZGXZ2UUK6X	Catherine Hallberg " (Kate)"	1	1
150509	0006641040	A3CMRKGE0P909G	Teresa	3	4



In [15]:

```
final = final[final.ProductId != "0006641040"]
```

In [16]:

```
print("Percentage of data still remaining : ",(final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100)
```

```
#Before starting the next phase of preprocessing Lets see the number of entries left  
print("Number of reviews left after Data Cleaning and Deduplication :")  
print(final.shape)
```

```
#How many positive and negative reviews are present in our dataset?  
final['Score'].value_counts()
```

```
Percentage of data still remaining : 69.14973735959865  
Number of reviews left after Data Cleaning and Deduplication :  
(363599, 10)
```

Out[16]:

```
positive    306566  
negative     57033  
Name: Score, dtype: int64
```

#### **Observation:-**

It is an imbalanced dataset as the number of positive reviews are way high in number than negative reviews.

## **[5] Text Preprocessing Using NLTK :**

In the Preprocessing phase we do the following in the order below:-

1. Removal of HTML Tags
2. Removal of any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Removal of Stopwords
7. Finally Snowball Stemming the word

After which we collect the words used to describe positive and negative reviews

In [20]:

```
# find sentences containing HTML tags
import re
i=0;
for sent in final['Text'].values:
    if (len(re.findall('<.*?>', sent))):
        print(i)
        print(sent)
        break;
    i += 1;
```

10

I wanted a treat that was accepted and well liked for my rescue animals.<br />This is the only treat that is healthy and loved by all 4 legged beings in my home!<br />It does not contain sugar or grains or silly vegetables which virtually all treats contain. Dogs, cats and ferrets are carnivores they are not cattle to eat grain or rabbits to eat vegetables, and WHYYYY do companies add sugar,beet pulp or corn syrup to carnivore foods? It is dangerous and can cause the death of an animal with diabetes.<br />It is pretty easy to break into smaller pieces for cats and kittens with weak jaws and its wonderful to use as an aid to gain the trust of an abused dog as it will not cause stomach upset when given in common sense amounts.<br />I like that it goes a long way as it costs alot to heal and maintain and train abused and rescued dogs.<br />NO minus to this product other then the price,I can not afford to use it as much as I would like.

## [5.1] Removal of html Tags :

In [21]:

```
#function to clean the word of any html-tags
def cleanhtml(sentence):
    cleanr = re.compile('<.*?>')
    cleantext = re.sub(cleanr, ' ', sentence)
    return cleantext
```

```
cleanhtml("<br />This is the only treat that is healthy and loved by all 4 legged beings in my home!<br />It does not contain sugar or grains")
```

Out[21]:

```
' This is the only treat that is healthy and loved by all 4 legged beings in my home! It does not contain sugar or grains'
```

## [5.2] Removal of Punctuations and unnecessary characters :

In [22]:

```
#function to clean the word of any punctuation or special characters
def cleanpunc(sentence):
    cleaned = re.sub(r'[?|!|\\"|#|@|~|%|*]',r'',sentence)
    cleaned = re.sub(r'[,|,|)|(|\|/]',r' ',cleaned)
    return cleaned

cleanpunc("WHYYYY do companies add sugar,beet pulp or corn syrup to carnivore foods?")
```

Out[22]:

```
'WHYYYY do companies add sugar beet pulp or corn syrup to carnivore foods'
```

## [5.3] StopWords :

In [23]:

```
import nltk
nltk.download('punkt')
nltk.download('stopwords')

stop = set(stopwords.words('english')) #set of stopwords
print(stop)
print("*****")
print("No. of stop words: ",len(stop))
```

```
[nltk_data] Downloading package punkt to /home/jovyan/nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package stopwords to /home/jovyan/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
{'further', 'should', 'didn', 'them', 'me', 'hasn't', 'as', 'ma', 'itsel
f', 'then', 'whom', 'weren't', 'has', 'from', 're', 'while', 'don't', 'are
n', 'haven't', 'but', 'such', 'very', 'didn't', 'nor', 'so', 'being', 'sha
n't', 'most', 'only', 'hadn', 'ourselves', 'because', 'what', 'were', 'is
n't', 'the', 'shouldn', 'should've', 'own', 'wouldn't', 'down', 'couldn',
'can', 'during', 'mightn', 'against', 'wasn't', 'up', 'hadn't', 'for', 'ot
her', 'hasn', 'myself', 'how', 'theirs', 'did', 'at', 'or', 'do', 't', 'ju
st', 'which', 'mightn't', 'that'll', 'if', 'we', 'this', 'after', 'now',
'mustn't', 'it's', 'they', 'than', 'hers', 'his', 'through', 'weren', 'ai
n', 'their', 'out', 'shan', 'won', 'all', 'had', 'are', 'those', 'above',
'you've', 'ours', 'between', 'she', 'our', 'haven', 'about', 'y', 'been',
'again', 'i', 'wasn', 'your', 'where', 'any', 'too', 'more', 'some', 'tha
t', 'her', 'below', 'each', 'yourself', 'he', 'herself', 'you'd', 'when',
'its', 'it', 's', 'once', 'needn't', 'there', 'themselves', 'of', 'off',
'was', 'with', 've', 'wouldn', 'she's', 'in', 'be', 'you're', 'few', 'no
t', 'into', 'yours', 'why', 'will', 'won't', 'on', 'doesn', 'having', 'is
n', 'to', 'both', 'does', 'my', 'am', 'couldn't', 'doesn't', 'no', 'you',
'who', 'have', 'doing', 'o', 'd', 'yourselves', 'him', 'these', 'is', 'und
er', 'here', 'aren't', 'by', 'a', 'shouldn't', 'himself', 'same', 'you'l
l', 'an', 'needn', 'and', 'over', 'll', 'until', 'm', 'don', 'before', 'mu
stn'}
*****
No. of stop words: 179
```

In [24]:

```
exceptions = ["aren't", "mightn", "wasn", "hadn", "don't", "against", "hadn't", "shan", "were  
n't", "didn", "don", "hasn't", \  
             "hasn", "shouldn", "didn't", "wouldn", "wasn't", "needn't", "shouldn't", "would  
n't", "aren", "isn't", "doesn't", \  
             "nor", "not", "needn", "couldn't", "mightn't", "mustn", "mustn't", "ain", "sha  
n't", "haven", "won't", "couldn", "isn", \  
             "weren", "haven't", "no", "haven't"]  
  
new_stop = []  
for i in stop:  
    if i not in exceptions:  
        new_stop.append(i)  
  
print("No. of stop words after removing exceptions: ", len(new_stop))
```

No. of stop words after removing exceptions: 140

## [5.4] Stemming :

In [25]:

```
sno = nltk.stem.SnowballStemmer('english') #initialising the snowball stemmer  
  
print("Original word: beautiful" + "|" + "Stem word: " + sno.stem('beautiful'))  
print("Original word: beauty" + "|" + "Stem word: " + sno.stem('beauty'))  
print("Original word: loved" + "|" + "Stem word: " + sno.stem('loved'))  
print("Original word: loving" + "|" + "Stem word: " + sno.stem('loving'))
```

```
Original word: beautiful|Stem word: beauti  
Original word: beauty|Stem word: beauti  
Original word: loved|Stem word: love  
Original word: loving|Stem word: love
```

### Observation:-

We can see words like "beautiful" and "beauty" have their stem as "beauti", "loved" and "loving" have their stem as "love".

Hence it helps in reducing the dimensions by taking the root stem of words.

## [5.5] Implementing the preprocessing steps one by one on all the reviews of dataset :

In [22]:

```
i=0
str1=' '
final_string=[]
all_positive_words=[] # store words from +ve reviews here
all_negative_words=[] # store words from -ve reviews here.
s=''

for sent in final['Text'].values:
    filtered_sentence=[]
    #print(sent);
    sent=cleanhtml(sent) # remove HTML tags
    for w in sent.split():
        for cleaned_words in cleanpunc(w).split():
            if((cleaned_words.isalpha()) & (len(cleaned_words)>2)):
                if(cleaned_words.lower() not in new_stop):
                    s=(sno.stem(cleaned_words.lower())).encode('utf8')
                    filtered_sentence.append(s)
                    if (final['Score'].values)[i] == 'positive':
                        all_positive_words.append(s) #list of all words used to describ
e positive reviews
                    if(final['Score'].values)[i] == 'negative':
                        all_negative_words.append(s) #list of all words used to describ
e negative reviews reviews
                    else:
                        continue
                else:
                    continue
    str1 = b" ".join(filtered_sentence) #final string of cleaned words
    #print("*****")

    final_string.append(str1)
    i+=1
```

In [23]:

```
from nltk.probability import FreqDist
pdist = FreqDist(all_positive_words)
top_positive = pdist.most_common(20)
print("Top 20 Positive words ocuring frequently in reviews:")
top_positive
```

Top 20 Positive words ocuring frequently in reviews:

Out[23]:

```
[(b'not', 146568),
 (b'like', 139160),
 (b'tast', 128865),
 (b'good', 112601),
 (b'flavor', 109329),
 (b'love', 107172),
 (b'use', 103792),
 (b'great', 103670),
 (b'one', 96529),
 (b'product', 90912),
 (b'tri', 86683),
 (b'tea', 83699),
 (b'coffe', 78763),
 (b'make', 75004),
 (b'get', 71996),
 (b'food', 64539),
 (b'would', 55477),
 (b'time', 55184),
 (b'buy', 54137),
 (b'realli', 52657)]
```



In [24]:

```
ndist = FreqDist(all_negative_words)
top_negative = ndist.most_common(20)
print("Top 20 Negative words ocuring frequently in reviews:")
top_negative
```

Top 20 Negative words ocuring frequently in reviews:

Out[24]:

```
[(b'not', 54325),
 (b'tast', 34534),
 (b'like', 32271),
 (b'product', 28181),
 (b'one', 20544),
 (b'flavor', 19520),
 (b'would', 17947),
 (b'tri', 17718),
 (b'use', 15280),
 (b'good', 15024),
 (b'coffe', 14700),
 (b'get', 13775),
 (b'buy', 13742),
 (b'order', 12862),
 (b'food', 12720),
 (b'dont', 11865),
 (b'tea', 11646),
 (b'even', 11068),
 (b'box', 10833),
 (b'amazon', 10067)]
```

## **[5.6] Adding a new column of CleanedText which displays the data after pre-processing of the review :**

In [27]:

```
final['CleanedText']=final_string
final['CleanedText']=final['CleanedText'].str.decode("utf-8")
```

In [26]:

```
final[['Text', 'CleanedText']].head(10)
```

Out[26]:

	Text	CleanedText
476617	This product by Arche...	product archer farm b...
22621	Our dogs just love th...	dog love saw pet stor...
22620	My dogs loves this ch...	dog love chicken prod...
284375	This book is easy to ...	book easi read ingred...
157850	I have been feeding m...	feed greyhound treat ...
157849	This is one product t...	one product welsh ter...
157833	This is the ONLY dog ...	dog treat lhasa apso ...
157832	These liver treas are...	liver trea phenomen r...
157837	This was the only tre...	treat dog like obedi ...
157831	No waste , even if sh...	wast even day goe hun...

## [5.7] Using SQLite Table to load data after preprocessing of reviews :

In [ ]:

```
# store final result into an SQLite table for future.
conn = sqlite3.connect('final.sqlite')
c=conn.cursor()
conn.text_factory = str
final.to_sql('Reviews', conn, schema=None, if_exists='replace', index=True, index_label=None, chunksize=None, dtype=None)
```

In [6]:

```
# using the SQLite Table to read data.
conn = sqlite3.connect('final.sqlite')

final = pd.read_sql_query(""" SELECT * FROM Reviews """,conn)
```

In [7]:

```
#Listing out the number of positive and negative reviews  
final = final.reset_index(drop=True)  
final['Score'].value_counts()
```

Out[7]:

```
positive    306566  
negative     57033  
Name: Score, dtype: int64
```

In [8]:

```
(final['Score'].value_counts()/len(final['Score']))*100
```

Out[8]:

```
positive    84.314313  
negative    15.685687  
Name: Score, dtype: float64
```

## [6] Train and Test Split of Data :

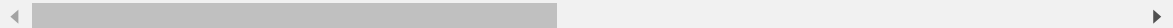
**Sorting the data by Time :**

In [10]:

```
final=final.sort_values('Time', axis=0, ascending=True, inplace=False, kind='quicksort', na_position='last')
final.head()
```

Out[10]:

	index	Id	ProductId	UserId	ProfileName	HelpfulnessNun
387	417839	451856	B00004CXX9	AIUWLEQ1ADEG5	Elizabeth Medina	0
293	346055	374359	B00004CI84	A344SMIA5JECGM	Vincent P. Ross	1
386	417838	451855	B00004CXX9	AJH6LUC1UT1ON	The Phantom of the Opera	0
209	346116	374422	B00004CI84	A1048CYU0OV4O8	Judy L. Eans	2
271	346041	374343	B00004CI84	A1B2IZU1JLZA6	Wes	19



### Time Based Slicing :

- Diving the data to Train set(first 70% ie older data) and Test Set(last 30% ie recent data)

In [11]:

```
from sklearn.model_selection import train_test_split

X = final["CleanedText"].values
y = final["Score"].values
X_train,X_test,y_train,y_test = train_test_split(X, y, test_size = 0.3,shuffle = False)
```

In [5]:

```
print("Shape of X_train: ",X_train.shape)
print("Shape of y_train: ",y_train.shape)
print("Shape of X_test: ",X_test.shape)
print("Shape of y_test: ",y_test.shape)
```

```
Shape of X_train: (254519,)
Shape of y_train: (254519,)
Shape of X_test: (109080,)
Shape of y_test: (109080,)
```

In [13]:

```
dumpfile(X,"X")
dumpfile(y,"y")
dumpfile(X_train,"X_train")
dumpfile(y_train,"y_train")
dumpfile(X_test,"X_test")
dumpfile(y_test,"y_test")
```

In [4]:

```
X = loadfile("X")
y = loadfile("y")
X_train = loadfile("X_train")
y_train = loadfile("y_train")
X_test = loadfile("X_test")
y_test = loadfile("y_test")
```

## [7] Naive Bayes :

### [7.1] Function to find the optimal alpha and error using K-fold cross-validation :

- Taking alpha between range 0.0001 and 1000.
- TimeSeries Split and performing K fold cross validation on Train Data
- Finding the optimal alpha
- Plotting between CV error/CV Accuracy and Alpha

**Bernoulli Naive Bayes** : BernoulliNB is suitable for discrete data designed for for binary/boolean features.

In [ ]:

```
from sklearn.model_selection import TimeSeriesSplit
from sklearn.naive_bayes import BernoulliNB
from sklearn.naive_bayes import MultinomialNB
from sklearn.model_selection import cross_val_score as cv
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.grid_search import GridSearchCV
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score
import warnings
warnings.filterwarnings('ignore')

alpha_values = [0.00001, 0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000]
param_grid = dict(alpha = alpha_values)

def BernoulliNB_Train(X_train, y_train):

    model = BernoulliNB()
    grid = GridSearchCV(model, param_grid, cv=10, scoring='accuracy')
    grid.fit(X_train, y_train)

    grid_mean_scores = [i.mean_validation_score for i in grid.grid_scores_]

    #Misclassification error
    MSE = [1 - x for x in grid_mean_scores]

    #Finding the optimal K
    optimal_alpha = grid.best_params_

    best_accuracy = np.round(grid.best_score_ * 100, 3)

    print("\n\033[1mOptimal alpha:\033[0m ", optimal_alpha)
    print("\n\033[1mCrossValidation Error:\033[0m {}".format(np.round(min(MSE), 3)))
    print("\n\033[1mCrossValidation Accuracy:\033[0m {} %\n\n".format(best_accuracy))

    plt.figure(figsize=(20, 6))
    plt.style.use('fivethirtyeight')
    plt.subplot(121)
    plt.plot(alpha_values, MSE, 'r-o')
    for xy in zip(alpha_values, np.round(MSE, 3)):
        plt.annotate('%s %s' % xy, xy = xy, textcoords = 'data')
    plt.title("CV Error vs Alpha Value")
    plt.xlabel("Alpha")
    plt.ylabel("CV Error")
    plt.grid(True)

    plt.subplot(122)
    plt.plot(alpha_values, grid_mean_scores, 'g-o')
    for xy in zip(alpha_values, np.round(grid_mean_scores, 3)):
        plt.annotate('%s %s' % xy, xy = xy, textcoords = 'data')
    plt.title("CV Accuracy vs Alpha Value")
    plt.xlabel("Alpha")
    plt.ylabel("CV Accuracy")
    plt.grid(True)
    plt.show()
```

```

print("\n\033[1mCV Error for each value of alpha:\033[0m ",np.round(MSE,3))
print("\n\033[1mCV Accuracy for each value of alpha:\033[0m ",np.round(grid_mean_scores,3))

```

**Multinomial Naive Bayes** :MultinomialNB is suitable for classification with discrete features(count based occurrences).

In [13]:

```

def MultinomialNB_Train(X_train,y_train):

    model = MultinomialNB()
    grid = GridSearchCV(model, param_grid, cv=10, scoring='accuracy')
    grid.fit(X_train, y_train)

    grid_mean_scores = [i.mean_validation_score for i in grid.grid_scores_]

    #Misclassification error
    MSE = [1 - x for x in grid_mean_scores]

    #Finding the optimal K
    optimal_alpha = grid.best_params_

    best_accuracy = np.round(grid.best_score_ * 100,3)

    print("\n\033[1mOptimal alpha:\033[0m ", optimal_alpha)
    print("\n\033[1mCrossValidation Error:\033[0m {}".format(np.round(min(MSE),3)))
    print("\n\033[1mCrossValidation Accuracy:\033[0m {} %\n\n".format(best_accuracy))

    plt.figure(figsize=(20,6))
    plt.style.use('fivethirtyeight')
    plt.subplot(121)
    plt.plot(alpha_values,MSE, 'r-o')
    for xy in zip(alpha_values, np.round(MSE,3)):
        plt.annotate('(%s %s)' % xy, xy = xy, textcoords = 'data')
    plt.title("CV Error vs Alpha Value")
    plt.xlabel("Alpha")
    plt.ylabel("CV Error")
    plt.grid(True)

    plt.subplot(122)
    plt.plot(alpha_values,grid_mean_scores, 'g-o')
    for xy in zip(alpha_values, np.round(grid_mean_scores,3)):
        plt.annotate('(%s %s)' % xy, xy = xy, textcoords = 'data')
    plt.title("CV Accuracy vs Alpha Value")
    plt.xlabel("Alpha")
    plt.ylabel("CV Accuracy")
    plt.grid(True)
    plt.show()

    print("\n\033[1mCV Error for each value of alpha:\033[0m ",np.round(MSE,3))
    print("\n\033[1mCV Accuracy for each value of alpha:\033[0m ",np.round(grid_mean_scores,3))

```

## [7.2] Function to find the features importance and predict on Test Data :

- Finding the features importance for each class
- Plotting the Confusion matrix
- Finding Accuracy, Precision, Recall and F1 Score on Test Data

**True Negative** : Number of datapoints with class label "negative" correctly classified as "negative".

**False Positive** : Number of datapoints with class label "negative" misclassified as "positive".

**False Negative** : Number of datapoints with class label "positive" misclassified as "negative".

**True Positive**: Number of datapoints with class label "positive" correctly classified as "positive".

**Precision** :  $\text{True Positive} / (\text{True Positive} + \text{False Positive})$

**Recall** :  $\text{True Positive} / (\text{True Positive} + \text{False Negative})$

**Bernoulli Naive Bayes** :



In [14]:

```
def BernoulliNB_Test(X_train,X_test,y_train,y_test,optimal_alpha,vectorizer,n = 25):

    optimal_model = BernoulliNB(alpha = optimal_alpha)
    optimal_model.fit(X_train, y_train)
    y_pred = optimal_model.predict(X_test)

    ##-----Feature Importance-----##
    class_labels = optimal_model.classes_
    feature_names = vectorizer.get_feature_names()
    top_negative = sorted(zip(optimal_model.coef_[0], feature_names))[:n]
    top_positive= sorted(zip(optimal_model.coef_[0], feature_names))[-n:]

    print("\033[1m-----Top {} Negative Words with high Importance-----\033[0m".format(n))
    neg_featureimp_df = pd.DataFrame(top_negative, columns=['Coefficient Factor','Features'])
    print(neg_featureimp_df.to_string(index=False))

    print("\n\033[1m-----Top {} Positive Words with high Importance-----\033[0m".format(n))
    pos_featureimp_df = pd.DataFrame(top_positive, columns=['Coefficient Factor','Features'],)
    print(pos_featureimp_df.to_string(index=False))

    print("\n\n")

    ##-----Confusion Matrix and Performance metrics
    accuracy = accuracy_score(y_test,y_pred) * 100
    precision = precision_score(y_test,y_pred,average= 'macro')
    recall = recall_score(y_test,y_pred,average= 'macro')
    f1= f1_score(y_test,y_pred,average= 'macro')
    MSE = (1 - (accuracy/100))
    cm = confusion_matrix(y_test, y_pred)
    tn, fp, fn, tp = confusion_matrix(y_test, y_pred).ravel()
    cm_df = pd.DataFrame(cm,
                        index = ['negative','positive'],
                        columns = ['negative','positive'])
    sns.heatmap(cm_df, annot=True)
    plt.title('Confusion Matrix')
    plt.ylabel('Actual Label')
    plt.xlabel('Predicted Label')
    plt.show()

    print(cm)
    print("\n\033[1mTest Error :\033[0m {}".format(np.round(MSE,3)))
    print("\033[1mTest Accuracy :\033[0m {} %".format(np.round(accuracy,3)))
    print("\033[1mTrue Negative :\033[0m {}".format(tn))
    print("\033[1mFalse Positive :\033[0m {}".format(fp))
    print("\033[1mFalse Negative :\033[0m {}".format(fn))
    print("\033[1mTrue Positive :\033[0m {}".format(tp))
    print("\033[1mPrecision Score :\033[0m {}".format(np.round(precision,3)))
    print("\033[1mRecall Score :\033[0m {}".format(np.round(recall,3)))
    print("\033[1mF1 Score :\033[0m {}".format(np.round(f1,3)))
```

---

**Multinomial Naive Bayes :**

In [15]:

```
def MultinomialNB_Test(X_train,X_test,y_train,y_test,optimal_alpha,vectorizer,n = 25):

    optimal_model = MultinomialNB(alpha = optimal_alpha)
    optimal_model.fit(X_train, y_train)
    y_pred = optimal_model.predict(X_test)

    ##-----Feature Importance-----##
    class_labels = optimal_model.classes_
    feature_names = vectorizer.get_feature_names()
    top_negative = sorted(zip(optimal_model.coef_[0], feature_names))[:n]
    top_positive= sorted(zip(optimal_model.coef_[0], feature_names))[-n:]

    print("\033[1m-----Top {} Negative Words with high Importance-----\033[0m".format(n))
    neg_featureimp_df = pd.DataFrame(top_negative, columns=['Coefficient Factor','Features'])
    print(neg_featureimp_df.to_string(index=False))

    print("\n\033[1m-----Top {} Positive Words with high Importance-----\033[0m".format(n))
    pos_featureimp_df = pd.DataFrame(top_positive, columns=['Coefficient Factor','Features'],)
    print(pos_featureimp_df.to_string(index=False))

    print("\n\n")

    ##-----Confusion Matrix and Performance metrics
    accuracy = accuracy_score(y_test,y_pred) * 100
    precision = precision_score(y_test,y_pred,average= 'macro')
    recall = recall_score(y_test,y_pred,average= 'macro')
    f1= f1_score(y_test,y_pred,average= 'macro')
    MSE = (1 - (accuracy/100))
    cm = confusion_matrix(y_test, y_pred)
    tn, fp, fn, tp = confusion_matrix(y_test, y_pred).ravel()
    cm_df = pd.DataFrame(cm,
                        index = ['negative','positive'],
                        columns = ['negative','positive'])
    sns.heatmap(cm_df, annot=True)
    plt.title('Confusion Matrix')
    plt.ylabel('Actual Label')
    plt.xlabel('Predicted Label')
    plt.show()

    print(cm)
    print("\n\033[1mTest Error :\033[0m {}".format(np.round(MSE,3)))
    print("\033[1mTest Accuracy :\033[0m {} %".format(np.round(accuracy,3)))
    print("\033[1mTrue Negative :\033[0m {}".format(tn))
    print("\033[1mFalse Positive :\033[0m {}".format(fp))
    print("\033[1mFalse Negative :\033[0m {}".format(fn))
    print("\033[1mTrue Positive :\033[0m {}".format(tp))
    print("\033[1mPrecision Score :\033[0m {}".format(np.round(precision,3)))
    print("\033[1mRecall Score :\033[0m {}".format(np.round(recall,3)))
    print("\033[1mF1 Score :\033[0m {}".format(np.round(f1,3)))
```

## [8] Featurization Methods :

**Note :** Naive Bayes is invariant to feature scaling.

### [8.1] Bag Of Words(unigram) :

**Count Based BOW :**

In [29]:

```
%%time
bow_unigram = CountVectorizer()
X_train_bowuni = bow_unigram.fit_transform(X_train)
print("Type of Count Vectorizer: ",type(X_train_bowuni))
print("The shape of text BOW vectorizer: ", X_train_bowuni.get_shape())
print("Number of unique word: ", X_train_bowuni.get_shape()[1])
```

```
Type of Count Vectorizer: <class 'scipy.sparse.csr.csr_matrix'>
The shape of text BOW vectorizer: (254519, 59315)
Number of unique word: 59315
CPU times: user 11.4 s, sys: 92 ms, total: 11.5 s
Wall time: 11.5 s
```

In [30]:

```
%%time
X_test_bowuni = bow_unigram.transform(X_test)
print("The shape of text BOW vectorizer: ", X_test_bowuni.get_shape())
print("Number of unique word: ", X_test_bowuni.get_shape()[1])
```

```
The shape of text BOW vectorizer: (109080, 59315)
Number of unique word: 59315
CPU times: user 5.42 s, sys: 8 ms, total: 5.43 s
Wall time: 5.43 s
```

In [31]:

```
dumpfile(X_train_bowuni,"X_train_bowuni")
dumpfile(X_test_bowuni,"X_test_bowuni")
```

In [85]:

```
X_train_bowuni = loadfile("X_train_bowuni")
X_test_bowuni = loadfile("X_test_bowuni")
```

In [35]:

```
print("Shape of Training Data: ",X_train_bowuni.shape)
print("Shape of Test Data: ",X_test_bowuni.shape)
```

```
Shape of Training Data: (254519, 59315)
Shape of Test Data: (109080, 59315)
```

**Binary BOW :**

In [8]:

```
%%time
binarybow_unigram = CountVectorizer(binary = True)
X_train_binarybowuni = binarybow_unigram.fit_transform(X_train)
print("Type of Count Vectorizer: ",type(X_train_binarybowuni))
print("The shape of text BOW vectorizer: ", X_train_binarybowuni.get_shape())
print("Number of unique word: ", X_train_binarybowuni.get_shape()[1])
```

```
Type of Count Vectorizer: <class 'scipy.sparse.csr.csr_matrix'>
The shape of text BOW vectorizer: (254519, 59315)
Number of unique word: 59315
CPU times: user 11.5 s, sys: 76 ms, total: 11.6 s
Wall time: 11.6 s
```

In [9]:

```
%%time
X_test_binarybowuni = binarybow_unigram.transform(X_test)
print("The shape of text BOW vectorizer: ", X_test_binarybowuni.get_shape())
print("Number of unique word: ", X_test_binarybowuni.get_shape()[1])
```

```
The shape of text BOW vectorizer: (109080, 59315)
Number of unique word: 59315
CPU times: user 5.4 s, sys: 12 ms, total: 5.41 s
Wall time: 5.41 s
```

In [10]:

```
print("Shape of Training Data: ",X_train_binarybowuni.shape)
print("Shape of Test Data: ",X_test_binarybowuni.shape)
```

```
Shape of Training Data: (254519, 59315)
Shape of Test Data: (109080, 59315)
```

**Using Bernoulli Naive Bayes(binary features) :**


In [16]:

```
%%time
if __name__ == "__main__":
    BernoulliNB_Train(X_train_binarybowuni, y_train)
```

**Optimal alpha:** {'alpha': 0.01}

**CrossValidation Error:** 0.11

**CrossValidation Accuracy:** 89.049 %



**CV Error for each value of alpha:** [0.111 0.11 0.11 0.11 0.111 0.116 0.145 0.149 0.149]

**CV Accuracy for each value of alpha:** [0.889 0.89 0.89 0.89 0.889 0.884 0.855 0.851 0.851]

CPU times: user 3min 2s, sys: 2.99 s, total: 3min 5s

Wall time: 3min 4s

In [19]:

```
%%time
if __name__ == "__main__":
    BernoulliNB_Test(X_train_binarybowuni, X_test_binarybowuni, y_train, y_test, 0.01,
        binarybow_unigram)
```


-----Top 25 Negative Words with high Importance-----

Coeficient	Factor	Features
-16.890669	aaaaaaarrrrrrggghhh	
-16.890669	aaaaaahhhhhhyaaaaaa	
-16.890669		aachen
-16.890669		aarrgh
-16.890669		aauc
-16.890669		abalon
-16.890669		abbazabba
-16.890669		abiet
-16.890669		abolitionist
-16.890669		abort
-16.890669		abottl
-16.890669		abrevi
-16.890669		abrotanum
-16.890669		absolutelt
-16.890669		absoprt
-16.890669		absurt
-16.890669		abswer
-16.890669		abvious
-16.890669		accepert
-16.890669		acceptal
-16.890669		acceptalbl
-16.890669		accor
-16.890669		accordng
-16.890669		accourd
-16.890669		accpet

-----Top 25 Positive Words with high Importance-----

Coeficient	Factor	Features
-2.078467		dont
-2.050982		eat
-2.040586		also
-2.032481		much
-2.014801		price
-1.984578		realli
-1.978280		find
-1.957953		best
-1.957887		would
-1.947343		amazon
-1.909697		time
-1.905528		buy
-1.695963		get
-1.646235		make
-1.563862		product
-1.533892		tri
-1.490217		use
-1.472885		one
-1.429562		flavor
-1.294476		great
-1.277035		good
-1.273151		love
-1.207503		tast
-1.190105		like
-1.143144		not





```
[[12567  6480]
 [ 6226 83807]]
```

**Test Error : 0.116**

**Test Accuracy : 88.352 %**

**True Negative : 12567**

**False Positive : 6480**

**False Negative : 6226**

**True Positive : 83807**

**Precision Score : 0.798**

**Recall Score : 0.795**

**F1 Score : 0.797**

CPU times: user 8.12 s, sys: 16 ms, total: 8.14 s

Wall time: 7.89 s

**Using Multinomial Naive Bayes :**


In [62]:

```
%%time
if __name__ == "__main__":
    MultinomialNB_Train(X_train_bowuni, y_train)
```

Optimal alpha: {'alpha': 1}

CrossValidation Error: 0.097

CrossValidation Accuracy: 90.299 %



CV Error for each value of alpha: [0.106 0.105 0.104 0.102 0.101 0.097 0.119 0.149 0.149]

CV Accuracy for each value of alpha: [0.894 0.895 0.896 0.898 0.899 0.903 0.881 0.851 0.851]

CPU times: user 2min 56s, sys: 1.7 s, total: 2min 57s  
Wall time: 2min 57s

In [87]:


```
%%time
if __name__ == "__main__":
    MultinomialNB_Test(X_train_bowuni, X_test_bowuni, y_train, y_test, 1, bow_unigram)
```

-----Top 25 Negative Words with high Importance-----

Coeficient	Factor	Features
-15.931921	aaaaaaarrrrrrggghhh	
-15.931921	aaaaaahhhhhhyaaaaaa	
-15.931921		aachen
-15.931921		aarrgh
-15.931921		aauc
-15.931921		abalon
-15.931921		abbazabba
-15.931921		abiet
-15.931921		abolitionist
-15.931921		abort
-15.931921		abottl
-15.931921		abrevi
-15.931921		abrotanum
-15.931921		absolutelt
-15.931921		absoprt
-15.931921		absurt
-15.931921		abswer
-15.931921		abvious
-15.931921		accepet
-15.931921		acceptal
-15.931921		acceptalbl
-15.931921		accor
-15.931921		accordng
-15.931921		accourd
-15.931921		accpet

-----Top 25 Positive Words with high Importance-----

Coeficient	Factor	Features
-5.482743		price
-5.482192		best
-5.466620		find
-5.429267		realli
-5.415981		eat
-5.412275		amazon
-5.394744		buy
-5.391936		would
-5.372002		time
-5.232076		food
-5.105465		get
-5.042244		make
-5.036089		coffe
-4.924601		tri
-4.904170		tea
-4.896048		product
-4.810499		one
-4.752973		use
-4.732351		great
-4.711891		love
-4.684199		flavor
-4.661488		good
-4.526458		tast
-4.456843		like
-4.397420		not



```
[[13353  5694]
 [ 5404 84629]]
```

**Test Error :** 0.102  
**Test Accuracy :** 89.826 %  
**True Negative :** 13353  
**False Positive :** 5694  
**False Negative :** 5404  
**True Positive :** 84629  
**Precision Score :** 0.824  
**Recall Score :** 0.821  
**F1 Score :** 0.822  
 CPU times: user 7.99 s, sys: 16 ms, total: 8.01 s  
 Wall time: 7.74 s

## [8.2] Bag Of Words(bigram) :

### Count Based BOW :

In [36]:

```
%%time
bow_bigram = CountVectorizer(ngram_range=(1, 2))
X_train_bowbi = bow_bigram.fit_transform(X_train)
print("Type of Count Vectorizer: ", type(X_train_bowbi))
print("The shape of text BOW vectorizer: ", X_train_bowbi.get_shape())
print("Number of unique word: ", X_train_bowbi.get_shape()[1])
```

Type of Count Vectorizer: <class 'scipy.sparse.csr.csr\_matrix'>  
 The shape of text BOW vectorizer: (254519, 2295006)  
 Number of unique word: 2295006  
 CPU times: user 36.9 s, sys: 320 ms, total: 37.3 s  
 Wall time: 37.3 s

In [37]:

```
%%time
X_test_bowbi = bow_bigram.transform(X_test)
print("The shape of text BOW vectorizer: ", X_test_bowbi.get_shape())
print("Number of unique word: ", X_test_bowbi.get_shape()[1])
```

The shape of text BOW vectorizer: (109080, 2295006)  
Number of unique word: 2295006  
CPU times: user 11.9 s, sys: 8 ms, total: 11.9 s  
Wall time: 11.9 s

In [38]:

```
dumpfile(X_train_bowbi,"X_train_bowbi")
dumpfile(X_test_bowbi,"X_test_bowbi")
```

In [88]:

```
X_train_bowbi = loadfile("X_train_bowbi")
X_test_bowbi = loadfile("X_test_bowbi")
```

In [42]:

```
print("Shape of Training Data: ",X_train_bowbi.shape)
print("Shape of Test Data: ",X_test_bowbi.shape)
```

Shape of Training Data: (254519, 2295006)  
Shape of Test Data: (109080, 2295006)

## Binary BOW :

In [20]:

```
%%time
binarybow_bigram = CountVectorizer(ngram_range=(1, 2), binary = True)
X_train_binarybowbi = binarybow_bigram.fit_transform(X_train)
print("Type of Count Vectorizer: ",type(X_train_binarybowbi))
print("The shape of text BOW vectorizer: ", X_train_binarybowbi.get_shape())
print("Number of unique word: ", X_train_binarybowbi.get_shape()[1])
```

Type of Count Vectorizer: <class 'scipy.sparse.csr.csr\_matrix'>  
The shape of text BOW vectorizer: (254519, 2295006)  
Number of unique word: 2295006  
CPU times: user 37.3 s, sys: 392 ms, total: 37.7 s  
Wall time: 37.7 s

In [21]:

```
%%time
X_test_binarybowbi = binarybow_bigram.transform(X_test)
print("The shape of text BOW vectorizer: ", X_test_binarybowbi.get_shape())
print("Number of unique word: ", X_test_binarybowbi.get_shape()[1])
```

The shape of text BOW vectorizer: (109080, 2295006)  
Number of unique word: 2295006  
CPU times: user 12.1 s, sys: 40 ms, total: 12.2 s  
Wall time: 12.2 s

In [22]:

```
print("Shape of Training Data: ",X_train_binarybowbi.shape)
print("Shape of Test Data: ",X_test_binarybowbi.shape)
```

Shape of Training Data: (254519, 2295006)

Shape of Test Data: (109080, 2295006)

## Using Bernoulli Naive Bayes(binary features) :


In [23]:

```
%%time
if __name__ == "__main__":
    BernoulliNB_Train(X_train_binarybowbi, y_train)
```

**Optimal alpha:** {'alpha': 0.001}

**CrossValidation Error:** 0.094

**CrossValidation Accuracy:** 90.626 %



**CV Error for each value of alpha:** [0.109 0.101 0.094 0.094 0.125 0.157 0.15 0.149 0.149]

**CV Accuracy for each value of alpha:** [0.891 0.899 0.906 0.906 0.875 0.843 0.85 0.851 0.851]

CPU times: user 4min 44s, sys: 10.8 s, total: 4min 55s

Wall time: 4min 55s

In [24]:

```
%%time
if __name__ == "__main__":
    BernoulliNB_Test(X_train_binarybowbi, X_test_binarybowbi, y_train, y_test, 0.001, binarybow_bigram)
```




-----Top 25 Negative Words with high Importance-----

Coeficient	Factor	Features
-19.193254		aaa dont
-19.193254		aaaaaaaarrrrrggghhh
-19.193254		aaaaaaaarrrrrggghhh back
-19.193254		aaaaaahhhhyaaaaaa
-19.193254		aaaaaahhhhyaaaaaa fire
-19.193254		aachen
-19.193254		aachen munich
-19.193254		aachen printen
-19.193254		aafco certifi
-19.193254		aafco definit
-19.193254		aafco regul
-19.193254		aamzon howev
-19.193254		aarrgh
-19.193254		aarrgh final
-19.193254		aauc
-19.193254		aauc shelv
-19.193254		aback flavor
-19.193254		aback foreign
-19.193254		aback main
-19.193254		aback potenc
-19.193254		aback presenc
-19.193254		aback smell
-19.193254		abalon
-19.193254		abalon like
-19.193254		abalon not

-----Top 25 Positive Words with high Importance-----

Coeficient	Factor	Features
-2.078468		dont
-2.050982		eat
-2.040586		also
-2.032481		much
-2.014801		price
-1.984578		realli
-1.978280		find
-1.957953		best
-1.957887		would
-1.947343		amazon
-1.909697		time
-1.905528		buy
-1.695964		get
-1.646235		make
-1.563862		product
-1.533892		tri
-1.490218		use
-1.472886		one
-1.429562		flavor
-1.294476		great
-1.277035		good
-1.273151		love
-1.207503		tast
-1.190105		like
-1.143144		not



```
[[ 7957 11090]  
 [ 1077 88956]]
```

Test Error : 0.112  
Test Accuracy : 88.846 %  
True Negative : 7957  
False Positive : 11090  
False Negative : 1077  
True Positive : 88956  
Precision Score : 0.885  
Recall Score : 0.703  
F1 Score : 0.751  
CPU times: user 20.6 s, sys: 128 ms, total: 20.7 s  
Wall time: 20.4 s

**Using Multinomial Naive Bayes :**


In [64]:

```
%%time
if __name__ == "__main__":
    MultinomialNB_Train(X_train_bowbi, y_train)
```

Optimal alpha: {'alpha': 1}

CrossValidation Error: 0.093

CrossValidation Accuracy: 90.726 %



CV Error for each value of alpha: [0.11 0.103 0.096 0.096 0.116 0.093 0.149 0.149 0.149]

CV Accuracy for each value of alpha: [0.89 0.897 0.904 0.904 0.884 0.907 0.851 0.851 0.851]

CPU times: user 4min 5s, sys: 7.76 s, total: 4min 13s

Wall time: 4min 13s

In [91]:


```
%%time
if __name__ == "__main__":
    MultinomialNB_Test(X_train_bowbi, X_test_bowbi, y_train, y_test, 1, bow_bigram)
```

-----Top 25 Negative Words with high Importance-----

Coeficient	Factor	Features
-16.73665		aaa dont
-16.73665		aaaaaaaarrrrrggghhh
-16.73665		aaaaaaaarrrrrggghhh back
-16.73665		aaaaaahhhhyaaaaaa
-16.73665		aaaaaahhhhyaaaaaa fire
-16.73665		aachen
-16.73665		aachen munich
-16.73665		aachen printen
-16.73665		aafco certifi
-16.73665		aafco definit
-16.73665		aafco regul
-16.73665		aamzon howev
-16.73665		aarrgh
-16.73665		aarrgh final
-16.73665		aauc
-16.73665		aauc shelv
-16.73665		aback flavor
-16.73665		aback foreign
-16.73665		aback main
-16.73665		aback potenc
-16.73665		aback presenc
-16.73665		aback smell
-16.73665		abalon
-16.73665		abalon like
-16.73665		abalon not

-----Top 25 Positive Words with high Importance-----

Coeficient	Factor	Features
-6.287471		price
-6.286921		best
-6.271349		find
-6.233996		realli
-6.220710		eat
-6.217004		amazon
-6.199473		buy
-6.196664		would
-6.176731		time
-6.036804		food
-5.910194		get
-5.846972		make
-5.840818		coffe
-5.729330		tri
-5.708898		tea
-5.700776		product
-5.615227		one
-5.557701		use
-5.537080		great
-5.516620		love
-5.488927		flavor
-5.466216		good
-5.331187		tast
-5.261571		like
-5.202148		not



```
[[ 7344 11703]
 [  515 89518]]
```

**Test Error :** 0.112  
**Test Accuracy :** 88.799 %  
**True Negative :** 7344  
**False Positive :** 11703  
**False Negative :** 515  
**True Positive :** 89518  
**Precision Score :** 0.909  
**Recall Score :** 0.69  
**F1 Score :** 0.741  
 CPU times: user 20.1 s, sys: 44 ms, total: 20.1 s  
 Wall time: 19.8 s

### [8.3] TF-IDF(unigram) :

#### Count Based TF-IDF :

In [43]:

```
%%time
tfidf_unigram = TfidfVectorizer()
X_train_tfidfuni = tfidf_unigram.fit_transform(X_train)
print("Type of Count Vectorizer: ", type(X_train_tfidfuni))
print("The shape of text TFIDF vectorizer: ", X_train_tfidfuni.get_shape())
print("Number of unique word: ", X_train_tfidfuni.get_shape()[1])
```

Type of Count Vectorizer: <class 'scipy.sparse.csr.csr\_matrix'>  
 The shape of text TFIDF vectorizer: (254519, 59315)  
 Number of unique word: 59315  
 CPU times: user 11.6 s, sys: 64 ms, total: 11.6 s  
 Wall time: 11.7 s

In [44]:

```
%%time
X_test_tfidfuni = tfidf_unigram.transform(X_test)
print("The shape of text TFIDF vectorizer: ", X_test_tfidfuni.get_shape())
print("Number of unique word: ", X_test_tfidfuni.get_shape()[1])
```

The shape of text TFIDF vectorizer: (109080, 59315)  
Number of unique word: 59315  
CPU times: user 5.54 s, sys: 8 ms, total: 5.54 s  
Wall time: 5.55 s

In [45]:

```
dumpfile(X_train_tfidfuni,"X_train_tfidfuni")
dumpfile(X_test_tfidfuni,"X_test_tfidfuni")
```

In [92]:

```
X_train_tfidfuni = loadfile("X_train_tfidfuni")
X_test_tfidfuni = loadfile("X_test_tfidfuni")
```

In [49]:

```
print("Shape of Training Data: ",X_train_tfidfuni.shape)
print("Shape of Test Data: ",X_test_tfidfuni.shape)
```

Shape of Training Data: (254519, 59315)  
Shape of Test Data: (109080, 59315)

### Binary TF-IDF :

In [25]:

```
%%time
binarytfidf_unigram = TfidfVectorizer(binary = True)
X_train_binarytfidfuni = binarytfidf_unigram.fit_transform(X_train)
print("Type of Count Vectorizer: ",type(X_train_binarytfidfuni))
print("The shape of text TFIDF vectorizer: ", X_train_binarytfidfuni.get_shape())
print("Number of unique word: ", X_train_binarytfidfuni.get_shape()[1])
```

Type of Count Vectorizer: <class 'scipy.sparse.csr.csr\_matrix'>  
The shape of text TFIDF vectorizer: (254519, 59315)  
Number of unique word: 59315  
CPU times: user 11.8 s, sys: 92 ms, total: 11.9 s  
Wall time: 11.9 s

In [26]:

```
%%time
X_test_binarytfidfuni = binarytfidf_unigram.transform(X_test)
print("The shape of text TFIDF vectorizer: ", X_test_binarytfidfuni.get_shape())
print("Number of unique word: ", X_test_binarytfidfuni.get_shape()[1])
```

The shape of text TFIDF vectorizer: (109080, 59315)  
Number of unique word: 59315  
CPU times: user 5.53 s, sys: 4 ms, total: 5.54 s  
Wall time: 5.53 s

In [27]:

```
print("Shape of Training Data: ",X_train_binarytfidfuni.shape)
print("Shape of Test Data: ",X_test_binarytfidfuni.shape)
```

Shape of Training Data: (254519, 59315)

Shape of Test Data: (109080, 59315)

## Using Bernoulli Naive Bayes(binary features) :


In [28]:

```
%%time
if __name__ == "__main__":
    BernoulliNB_Train(X_train_binarytfidfuni, y_train)
```

**Optimal alpha:** {'alpha': 0.01}

**CrossValidation Error:** 0.11

**CrossValidation Accuracy:** 89.049 %



**CV Error for each value of alpha:** [0.111 0.11 0.11 0.11 0.111 0.116 0.145 0.149 0.149]

**CV Accuracy for each value of alpha:** [0.889 0.89 0.89 0.89 0.889 0.884 0.855 0.851 0.851]

CPU times: user 3min 2s, sys: 1.86 s, total: 3min 3s

Wall time: 3min 3s



In [29]:


```
%%time
if __name__ == "__main__":
    BernoulliNB_Test(X_train_binarytfidfuni, X_test_binarytfidfuni, y_train, y_test, 0.01, binarytfidf_unigram)
```

-----Top 25 Negative Words with high Importance-----

Coeficient	Factor	Features
-16.890669	aaaaaaarrrrrrggghhh	
-16.890669	aaaaaahhhhhhyaaaaaa	
-16.890669		aachen
-16.890669		aarrgh
-16.890669		aauc
-16.890669		abalon
-16.890669		abbazabba
-16.890669		abiet
-16.890669		abolitionist
-16.890669		abort
-16.890669		abottl
-16.890669		abrevi
-16.890669		abrotanum
-16.890669		absolutelt
-16.890669		absoprt
-16.890669		absurt
-16.890669		abswer
-16.890669		abvious
-16.890669		accepert
-16.890669		acceptal
-16.890669		acceptalbl
-16.890669		accor
-16.890669		accordng
-16.890669		accourd
-16.890669		accpet

-----Top 25 Positive Words with high Importance-----

Coeficient	Factor	Features
-2.078467		dont
-2.050982		eat
-2.040586		also
-2.032481		much
-2.014801		price
-1.984578		realli
-1.978280		find
-1.957953		best
-1.957887		would
-1.947343		amazon
-1.909697		time
-1.905528		buy
-1.695963		get
-1.646235		make
-1.563862		product
-1.533892		tri
-1.490217		use
-1.472885		one
-1.429562		flavor
-1.294476		great
-1.277035		good
-1.273151		love
-1.207503		tast
-1.190105		like
-1.143144		not



```
[[12567  6480]
 [ 6226 83807]]
```

Test Error : 0.116  
Test Accuracy : 88.352 %  
True Negative : 12567  
False Positive : 6480  
False Negative : 6226  
True Positive : 83807  
Precision Score : 0.798  
Recall Score : 0.795  
F1 Score : 0.797  
CPU times: user 8.12 s, sys: 12 ms, total: 8.13 s  
Wall time: 7.86 s

**Using Multinomial Naive Bayes :**


In [66]:

```
%%time
if __name__ == "__main__":
    MultinomialNB_Train(X_train_tfidfuni, y_train)
```

**Optimal alpha:** {'alpha': 0.01}

**CrossValidation Error:** 0.117

**CrossValidation Accuracy:** 88.26 %



**CV Error for each value of alpha:** [0.121 0.12 0.119 0.117 0.118 0.139 0.149 0.149 0.149]

**CV Accuracy for each value of alpha:** [0.879 0.88 0.881 0.883 0.882 0.861 0.851 0.851 0.851]

CPU times: user 2min 56s, sys: 1.32 s, total: 2min 58s

Wall time: 2min 58s

In [94]:


```
%%time
if __name__ == "__main__":
    MultinomialNB_Test(X_train_tfidfuni, X_test_tfidfuni, y_train, y_test, 0.01, tfidf_
unigram)
```

-----Top 25 Negative Words with high Importance-----

Coeficient	Factor	Features
-18.434172	aaaaaaarrrrrrggghhh	
-18.434172	aaaaaahhhhhhyaaaaaa	
-18.434172		aachen
-18.434172		aarrgh
-18.434172		aauc
-18.434172		abalon
-18.434172		abbazabba
-18.434172		abiet
-18.434172		abolitionist
-18.434172		abort
-18.434172		abottl
-18.434172		abrevi
-18.434172		abrotanum
-18.434172		absolutelt
-18.434172		absoprt
-18.434172		absurt
-18.434172		abswer
-18.434172		abvious
-18.434172		accepert
-18.434172		acceptal
-18.434172		acceptalbl
-18.434172		accor
-18.434172		accordng
-18.434172		accourd
-18.434172		accpet

-----Top 25 Positive Words with high Importance-----

Coeficient	Factor	Features
-5.781033		eat
-5.754349		realli
-5.747922		time
-5.742701		order
-5.710403		food
-5.709029		find
-5.691603		amazon
-5.671766		buy
-5.650831		best
-5.640605		price
-5.610394		get
-5.550292		make
-5.486820		tri
-5.419118		one
-5.336020		use
-5.320238		product
-5.249362		coffe
-5.220498		flavor
-5.205193		not
-5.179500		tea
-5.164250		like
-5.146301		good
-5.146180		tast
-5.083053		love
-5.069295		great



```
[[ 4952 14095]
 [  745 89288]]
```

**Test Error** : 0.136  
**Test Accuracy** : 86.395 %  
**True Negative** : 4952  
**False Positive** : 14095  
**False Negative** : 745  
**True Positive** : 89288  
**Precision Score** : 0.866  
**Recall Score** : 0.626  
**F1 Score** : 0.662  
 CPU times: user 8.11 s, sys: 20 ms, total: 8.13 s  
 Wall time: 7.87 s

## [8.4] TF-IDF(bigram) :

### Count Based TF-IDF :

In [50]:

```
%%time
tfidf_bigram = TfidfVectorizer(ngram_range=(1, 2))
X_train_tfidfbi = tfidf_bigram.fit_transform(X_train)
print("Type of Count Vectorizer: ", type(X_train_tfidfbi))
print("The shape of text TFIDF vectorizer: ", X_train_tfidfbi.get_shape())
print("Number of unique word: ", X_train_tfidfbi.get_shape()[1])
```

Type of Count Vectorizer: <class 'scipy.sparse.csr.csr\_matrix'>  
 The shape of text TFIDF vectorizer: (254519, 2295006)  
 Number of unique word: 2295006  
 CPU times: user 39.2 s, sys: 348 ms, total: 39.6 s  
 Wall time: 39.6 s

In [51]:

```
%%time
X_test_tfidfbi = tfidf_bigram.transform(X_test)
print("The shape of text TFIDF vectorizer: ", X_test_tfidfbi.get_shape())
print("Number of unique word: ", X_test_tfidfbi.get_shape()[1])
```

The shape of text TFIDF vectorizer: (109080, 2295006)  
Number of unique word: 2295006  
CPU times: user 12.7 s, sys: 32 ms, total: 12.7 s  
Wall time: 12.7 s

In [52]:

```
dumpfile(X_train_tfidfbi,"X_train_tfidfbi")
dumpfile(X_test_tfidfbi,"X_test_tfidfbi")
```

In [95]:

```
X_train_tfidfbi = loadfile("X_train_tfidfbi")
X_test_tfidfbi = loadfile("X_test_tfidfbi")
```

In [56]:

```
print("Shape of Training Data: ",X_train_tfidfbi.shape)
print("Shape of Test Data: ",X_test_tfidfbi.shape)
```

Shape of Training Data: (254519, 2295006)  
Shape of Test Data: (109080, 2295006)

### Binary TF-IDF :

In [30]:

```
%%time
binarytfidf_bigram = TfidfVectorizer(ngram_range=(1, 2), binary = True)
X_train_binarytfidfbi = binarytfidf_bigram.fit_transform(X_train)
print("Type of Count Vectorizer: ",type(X_train_binarytfidfbi))
print("The shape of text TFIDF vectorizer: ", X_train_binarytfidfbi.get_shape())
print("Number of unique word: ", X_train_binarytfidfbi.get_shape()[1])
```

Type of Count Vectorizer: <class 'scipy.sparse.csr.csr\_matrix'>  
The shape of text TFIDF vectorizer: (254519, 2295006)  
Number of unique word: 2295006  
CPU times: user 39.9 s, sys: 452 ms, total: 40.3 s  
Wall time: 40.3 s

In [31]:

```
%%time
X_test_binarytfidfbi = binarytfidf_bigram.transform(X_test)
print("The shape of text TFIDF vectorizer: ", X_test_binarytfidfbi.get_shape())
print("Number of unique word: ", X_test_binarytfidfbi.get_shape()[1])
```

The shape of text TFIDF vectorizer: (109080, 2295006)  
Number of unique word: 2295006  
CPU times: user 12.9 s, sys: 48 ms, total: 13 s  
Wall time: 13 s



In [32]:

```
print("Shape of Training Data: ",X_train_binarytfidfbi.shape)
print("Shape of Test Data: ",X_test_binarytfidfbi.shape)
```

Shape of Training Data: (254519, 2295006)

Shape of Test Data: (109080, 2295006)

## Using Bernoulli Naive Bayes(binary features) :


In [33]:

```
%%time
if __name__ == "__main__":
    BernoulliNB_Train(X_train_binarytfidfbi, y_train)
```

**Optimal alpha:** {'alpha': 0.001}

**CrossValidation Error:** 0.094

**CrossValidation Accuracy:** 90.626 %



**CV Error for each value of alpha:** [0.109 0.101 0.094 0.094 0.125 0.157 0.15 0.149 0.149]

**CV Accuracy for each value of alpha:** [0.891 0.899 0.906 0.906 0.875 0.843 0.85 0.851 0.851]

CPU times: user 4min 45s, sys: 10 s, total: 4min 55s

Wall time: 4min 55s

In [34]:


```
%%time
if __name__ == "__main__":
    BernoulliNB_Test(X_train_binarytfidfbi, X_test_binarytfidfbi, y_train, y_test, 0.00
1, binarytfidf_bigram)
```

-----Top 25 Negative Words with high Importance-----

Coeficient	Factor	Features
-19.193254		aaa dont
-19.193254		aaaaaaaarrrrrggghhh
-19.193254		aaaaaaaarrrrrggghhh back
-19.193254		aaaaaahhhhyaaaaaa
-19.193254		aaaaaahhhhyaaaaaa fire
-19.193254		aachen
-19.193254		aachen munich
-19.193254		aachen printen
-19.193254		aafco certifi
-19.193254		aafco definit
-19.193254		aafco regul
-19.193254		aamzon howev
-19.193254		aarrgh
-19.193254		aarrgh final
-19.193254		aauc
-19.193254		aauc shelv
-19.193254		aback flavor
-19.193254		aback foreign
-19.193254		aback main
-19.193254		aback potenc
-19.193254		aback presenc
-19.193254		aback smell
-19.193254		abalon
-19.193254		abalon like
-19.193254		abalon not

-----Top 25 Positive Words with high Importance-----

Coeficient	Factor	Features
-2.078468		dont
-2.050982		eat
-2.040586		also
-2.032481		much
-2.014801		price
-1.984578		realli
-1.978280		find
-1.957953		best
-1.957887		would
-1.947343		amazon
-1.909697		time
-1.905528		buy
-1.695964		get
-1.646235		make
-1.563862		product
-1.533892		tri
-1.490218		use
-1.472886		one
-1.429562		flavor
-1.294476		great
-1.277035		good
-1.273151		love
-1.207503		tast
-1.190105		like
-1.143144		not



```
[[ 7957 11090]  
 [ 1077 88956]]
```

Test Error : 0.112  
Test Accuracy : 88.846 %  
True Negative : 7957  
False Positive : 11090  
False Negative : 1077  
True Positive : 88956  
Precision Score : 0.885  
Recall Score : 0.703  
F1 Score : 0.751  
CPU times: user 21 s, sys: 100 ms, total: 21.1 s  
Wall time: 20.8 s

**Using Multinomial Naive Bayes :**


In [68]:

```
%%time
if __name__ == "__main__":
    MultinomialNB_Train(X_train_tfidfbi, y_train)
```

Optimal alpha: {'alpha': 0.001}

CrossValidation Error: 0.097

CrossValidation Accuracy: 90.265 %



CV Error for each value of alpha: [0.114 0.106 0.097 0.098 0.103 0.149 0.149 0.149 0.149]

CV Accuracy for each value of alpha: [0.886 0.894 0.903 0.902 0.897 0.851 0.851 0.851 0.851]

CPU times: user 4min 8s, sys: 6.21 s, total: 4min 14s

Wall time: 4min 14s

In [97]:


```
%%time
if __name__ == "__main__":
    MultinomialNB_Test(X_train_tfidfbi, X_test_tfidfbi, y_train, y_test, 0.001, tfidf_b
igram)
```

-----Top 25 Negative Words with high Importance-----

Coeficient	Factor	Features
-21.140633		aaa dont
-21.140633		aaaaaaaarrrrrggghhh
-21.140633		aaaaaaaarrrrrggghhh back
-21.140633		aaaaaahhhhyaaaaaa
-21.140633		aaaaaahhhhyaaaaaa fire
-21.140633		aachen
-21.140633		aachen munich
-21.140633		aachen printen
-21.140633		aafco certifi
-21.140633		aafco definit
-21.140633		aafco regul
-21.140633		aamzon howev
-21.140633		aarrgh
-21.140633		aarrgh final
-21.140633		aauc
-21.140633		aauc shelv
-21.140633		aback flavor
-21.140633		aback foreign
-21.140633		aback main
-21.140633		aback potenc
-21.140633		aback presenc
-21.140633		aback smell
-21.140633		abalon
-21.140633		abalon like
-21.140633		abalon not

-----Top 25 Positive Words with high Importance-----

Coeficient	Factor	Features
-6.913309		eat
-6.891037		realli
-6.890222		order
-6.884229		time
-6.848279		find
-6.834056		amazon
-6.820494		buy
-6.818411		food
-6.795146		price
-6.787423		best
-6.740259		get
-6.671618		make
-6.611024		tri
-6.546318		one
-6.464482		product
-6.454661		use
-6.374521		coffe
-6.351155		flavor
-6.327687		not
-6.290613		good
-6.288916		like
-6.283790		tea
-6.279446		tast
-6.231470		love
-6.222472		great



```
[[ 6870 12177]
 [  821 89212]]
```

Test Error : 0.119

Test Accuracy : 88.084 %

True Negative : 6870

False Positive : 12177

False Negative : 821

True Positive : 89212

Precision Score : 0.887

Recall Score : 0.676

F1 Score : 0.723

CPU times: user 29.6 s, sys: 100 ms, total: 29.7 s

Wall time: 29.4 s



## [9] Conclusion :

Featurization Model	Bernoulli NB				Multinomial Naive Bayes			
	Accuracy	Precision	Recall	F1 score	Accuracy	Precision	Recall	F1 score
BOW(unigram)	88.352 %	0.798	0.795	0.797	89.826 %	0.824	0.821	0.822
BOW(bigram)	88.846 %	0.885	0.703	0.751	88.799 %	0.909	0.69	0.741
TF-IDF(unigram)	88.352 %	0.798	0.795	0.797	86.395 %	0.866	0.626	0.661
TF-IDF(bigram)	88.846 %	0.885	0.703	0.751	88.084 %	0.887	0.676	0.723

**1** - Using **Bag Of Words(unigram)** method, **Multinomial Naive Bayes** model gives best accuracy of **89.826 %** and F1 score of **0.822** with **alpha = 1**.

**2**- Both MultinomialNB and BernouliNB are suitable for discrete data. Multinomial NB works with occurrence count of words while later is suitable for binary/boolean features.

**3** - It is also observed that Run Time complexity of Naive Bayes is **super fast** as compared to KNN. Thus, Naive Bayes can be used as a classification model for **low latency applications**.

**4** - Naive Bayes holds a concept of "**conditional independence**" that states features should be independent of each other. Hence featurization techniques like **average word2vec** and **tfidf weighted word2vec** which involves dependency between features **doesnot perform good** with Naive Bayes.